Multi-instance Publicly Verifiable Time-lock Puzzle and its Applications

Abstract. Time-lock puzzles are elegant protocols that enable a party to lock a message such that no one else can unlock it until a certain time elapses. Nevertheless, existing schemes are not suitable for the case where a server is given *multiple instances* of a puzzle scheme at once and it must unlock them at different points in time. If the schemes are naively used in this setting, then the server has to start solving all puzzles as soon as it receives them, that ultimately imposes significant computation cost and demands a high level of parallelisation. We put forth and formally define a primitive called "*multi-instance time-lock puzzle*" which allows composing a puzzle's instances. We propose a candidate construction: "*chained time-lock puzzle*" (C-TLP). It allows the server, given instances' composition, to solve puzzles sequentially, without having to run parallel computations on them. C-TLP makes black-box use of a standard time-lock puzzle scheme and is accompanied by a lightweight publicly verifiable algorithm. It is the first time-lock puzzle that offers a combination of the above features. We utilise C-TLP to build the first "outsourced proofs of retrievability" that can support *real-time detection* and *fair payment* while having lower overhead than the state of the art. As another application of C-TLP, we illustrate in certain cases, one can substitute a "verifiable delay function" with C-TLP, to gain much better efficiency.

1 Introduction

Time-lock puzzles are interesting cryptographic primitives that allow sending information to the future. They enable a party to lock a message such that, no one else can unlock it until a certain time has passed[1]. They have a wide range of applications, such as e-voting [16], fair contract signing [11], and sealed-bid auctions [43]. Over the last two decades, a variety of time-lock puzzles have been proposed. Nevertheless, existing puzzle schemes do not offer any efficient remedy for the multi-instance setting, where a server is given multiple instances of a puzzle at once and it should find one puzzle's solution after another. It is a natural generalisation of the single puzzle setting. Application areas include, but not limited to: (a) mass release of confidential documents over time, (b) gradually revealing multiple secret keys, (c) verifying continuous availability of cloud's services, e.g. data storage or secure hardware, or (d) scheduled private payments, where not only is every payment made after a certain period, but also the payment details remain confidential during the period. If existing puzzle schemes are utilised directly in the multi-instance setting, then the server has to deal with all puzzle instances, right after it receives them. This results in a significant computation overhead and requires a high level of parallelisation.

---

[1] There exist protocols that use an assistance of a third party to support time-release of a secret. This protocols' category is not our focus in this paper.

In this paper, we propose "*multi-instance time-luck puzzle*", a primitive that allows composing a puzzle's instances, where given the composition, a server can deal with each instance sequentially. We formally define the primitive and present an instantiation of it, "*chained time-lock puzzle*" (C-TLP). It makes black-box use of a standard time-lock puzzle scheme and is equipped with a *lightweight* verification algorithm that allows anyone to check the correctness of a solution found by the server. Its overall computation complexity of solving z puzzles is equivalent to that of solving only the last puzzle. The same procedure also imposes a communication overhead linear with z, i.e. $O(z)$. C-TLP is the first time-lock puzzle scheme that offers all the above features. Furthermore, we present concrete applications of the primitive and demonstrate its use case in a blockchain-based solution. Specifically, we combine the primitive's instantiation with a smart contract and apply the combination to "outsourced proofs of retrievability" research line, and propose "*smarter outsourced proofs of retrievability*" (SO-PoR) scheme which offers a combination of real-time detection and fair payment while imposing very low overhead, that makes it particularly suitable for mission-critical data. SO-PoR verification and store phases impose $\frac{1}{4.5}$ and $\frac{1}{46\times10^5}$ of computation costs imposed by the same phases in the fastest outsourced PoR. A server-side bandwidth of SO-PoR is much lower too; for instance, for a 1-GB file and 100 verifications, a server in SO-PoR requires 9×10^4 times fewer bits than those required in the state of the art protocol. Also, we show under certain circumstances C-TLP can play the role of a "verifiable delay function" (VDF) but with much lower overhead, i.e. a prover's computation and communication costs will be reduced by factors of 3 and 6.5 respectively.

Summary of Our Contributions. We (a) put forth the notion of multi-instance time-lock puzzle, formally define it, and identify its concrete applications, (b) present a candidate construction, C-TLP, the first multi-instance time-lock puzzle that is built on a standard time-lock puzzle, supports public verifiability, and has low costs, (c) propose the first outsourced PoR that can offer real-time detection and fair payment while maintaining low costs, and (d) show in certain cases, a VDF can be replaced with C-TLP to gain better efficiency.

2   Related Work

In this section, we provide a summary of related work. For a comprehensive survey, we refer readers to Appendix A.

Time-lock Puzzles The idea to send information into the *future*, i.e. time-lock puzzle/encryption, was first put forth by Timothy C. May. A time-lock puzzle allows a party to encrypt a message such that it cannot be decrypted until a certain time has passed. In general, a time-lock scheme should allow generating (and verifying) a puzzle to take less time than solving it. The scheme that May proposed lies on a trusted agent. Later, Rivest *et al.* [43] propose an RSA-based puzzle scheme that does not require a trusted agent, and is secure against a receiver who may have access to many computation resources that run in parallel. The latter protocol has been the core of (almost) all later time-lock puzzle schemes that supports the encapsulation of an arbitrary message. Later, [11, 21] proposed a scheme which also let a puzzle generator prove (in Zero-knowledge) to a puzzle solver that the correct solution will be recovered after a

certain time. Recently, [39, 12] propose homomorphic time-lock puzzles, where an arbitrary function can be run over puzzles before they are solved. In the protocols, all puzzles have an identical time parameter, and their solutions are supposed to be discovered at the same time. They are based on the RSA-based puzzle and fully homomorphic encryption which is computationally expensive.

Outsourced Proofs of Retrievability Proofs of retrievability (PoR) schemes, introduced in [26], guarantee to a client that its data on a cloud server is fully accessible. Ever since a variety of PoR's has been proposed. Recently, [3, 53] present *outsourced* PoR protocols that allow clients to outsource the verification to a potentially malicious third-party auditor. The scheme in [3] has the fastest prove and verification algorithms. It uses message authentication code (MAC) based tags, zero-knowledge proofs and error-correcting codes. However, it has several shortcomings, i.e. it offers no real-time detection, provides no efficient way for fair payments and has high costs of setup and auditor onboarding. Xu *et al.* in [53] propose a publicly verifiable outsourced PoR to improve the previous scheme's setup cost. It uses BLS signatures-like tags, polynomial arithmetics and error-correcting code. In this protocol, an auditor is assumed to be fully trusted during each verification whose overhead is higher than [3]. Very recently in [5] two protocols are proposed; namely, "basic PoSt" and "compact PoSt". They ensure that a client's data remains available on a server for a period, without the client's involvement in that period. The basic PoSt uses a Merkle tree-based PoR and VDF. It has a high communication cost. Since it requires a verifier to validate VDF's outputs, it imposes a significant computation cost too. The compact PoSt has a lower communication cost than the basic one, as it let the server combine PoR proofs. The protocol is mainly based on a trapdoor delay function (TDF). Note, all outsourced PoR schemes [3, 53, 5] assume the client behaves honestly towards the server. Otherwise, a malicious client can generate the tags in a way that makes an honest server generate invalid proofs.

Blockchain-based PoR. There exist distributed PoR schemes that let a client distribute its file among different (tailored) blockchain nodes, e.g. Permacoin [40], Filecoin [35], and KopperCoin [30]. However, they have either a large proof size (e.g. in [40, 35]) logarithmic with the file size when a Merkle tree is used, or high verification overhead (e.g. in [30]) due to the use of BLS signatures. There are protocols that use blockchain to verify the retrievability of off-chain data [42, 25, 57, 20, 8, 49]. Nevertheless, they either impose a high communication/computation cost [42, 25, 20, 8, 49], or clients have to be online for each verification [57]. Campanelli *et al.* [13] present a fair exchange mechanism over a blockchain that ensures the server gets paid if it provides an accepting PoR proof. But, this scheme assumes either the client can perform the verification itself or a third-party, acting on the client's behalf, carries out the verification honestly.

3 Preliminaries

In this section we provide the main primitives used in this work. We provide a notation table in Appendix B.

3.1 Smart Contract

Cryptocurrencies, such as Bitcoin and Ethereum, in addition to offering a decentralised currency, support computations on transactions. In this setting, often a certain com-

putation logic is encoded in a computer program, called *"smart contract"*. To date, Ethereum is the most predominant cryptocurrency framework that enables users to define arbitrary smart contracts. In this framework, contract code is stored on the blockchain and executed by all parties (i.e. miners) maintaining the cryptocurrency, when the program inputs are provided by transactions. The program execution's correctness is guaranteed by the security of the underlying blockchain components. To prevent a denial of service attack, the framework requires a transaction creator to pay a fee, called *"gas"*, depending on the complexity of the contract running on it.

3.2 Commitment Scheme

A commitment scheme involves two parties: *sender* and *receiver*, and includes two phases: *commit* and *open*. In the commit phase, the sender commits to a message: m as $\text{Com}(m, d) = h$, that involves a secret value: d. At the end of the commit phase, the commitment: h is sent to the receiver. In the open phase, the sender sends the opening: $\ddot{p} = (m, d)$ to the receiver who verifies its correctness: $\text{Ver}(h, \ddot{p}) \overset{?}{=} 1$ and accepts if the output is 1. A commitment scheme must satisfy two properties: (a) *hiding*: infeasible for an adversary (i.e. the receiver) to learn any information about the committed message: m, until the commitment: h is opened, and (b) *binding*: infeasible for an adversary (i.e. the sender) to open a commitment: h to different values: $\ddot{p}' = (m', d')$ than that used in the commit phase, i.e. infeasible to find $\ddot{p}'$, *s.t.* $\text{Ver}(h, \ddot{p}) = \text{Ver}(h, \ddot{p}') = 1$, where $\ddot{p} \neq \ddot{p}'$. There exist efficient non-interactive commitment schemes both in (a) the random oracle model using the well-known hash-based scheme such that $\text{Com}(m, d)$ involves computing: $\text{H}(m||d) = h$ and $\text{Ver}(h, \ddot{p})$ requires checking: $\text{H}(m||d) \overset{?}{=} h$, where H is a hash function, and (b) the standard model, e.g. Pedersen scheme [41].

3.3 Pseudorandom Function

Informally, a pseudorandom function (PRF) is a deterministic function that takes a key and an input; and outputs a value indistinguishable from that of a truly random function with the same input. A PRF is formally defined as follows [29].

Definition 1. *Let* $W : \{0, 1\}^{\psi} \times \{0, 1\}^{\eta} \to \{0, 1\}^{\iota}$ *be an efficient keyed function. It is said* W *is a pseudorandom function if for all probabilistic polynomial-time distinguishers* B, *there is a negligible function,* $\mu(.)$, *such that:*

$$\left| Pr[B^{W_k(\cdot)}(1^{\psi}) = 1] - Pr[B^{\omega(\cdot)}(1^{\psi}) = 1] \right| \leq \mu(\psi)$$

where the key, $k \overset{\$}{\leftarrow} \{0, 1\}^{\psi}$, *is chosen uniformly at random and* ω *is chosen uniformly at random from the set of functions mapping* η-*bit strings to* ι-*bit strings.*

3.4 Time-lock Puzzle

In this section, we restate the formal definition of a time-lock puzzle as well as RSA-based time-lock puzzle protocol [43]. We consider the RSA-based puzzle because of its simplicity and being the core of (almost) all later time-lock puzzle schemes.

Definition 2 (Time-lock Puzzle). *A time-lock puzzle comprises the following efficient three algorithms, such that the puzzle satisfies completeness and efficiency properties.*

- **Algorithms:**
  - Setup$(1^\lambda, \Delta) \to (pk, sk)$: *a probabilistic algorithm that takes an input security: 1^λ and time: Δ parameters. It outputs a public-private key pairs: (pk, sk)*
  - GenPuz$(s, pk, sk) \to \ddot{o}$: *a probabilistic algorithm that takes an input a solution: s and the public-private key pairs: (pk, sk). It outputs a puzzle: $\ddot{o}$*
  - SolvPuz$(pk, \ddot{o}) \to s$: *a deterministic algorithm that takes as input the public key: pk and puzzle: $\ddot{o}$. It outputs a solution: s*
- **Completeness**: *always* SolvPuz$(pk, \text{GenPuz}(s, pk, sk)) = s$
- **Efficiency**: *the run-time of algorithm* SolvPuz$(pk, \ddot{o})$ *is bounded by* $poly(\Delta, \lambda)$, *where* $poly(.)$ *is a polynomial.*

Informally, a time-lock puzzle's security requires that the puzzle solution remain hidden from all adversaries running in parallel within the time period, Δ. It is essential that no adversary can find a solution in time $\delta(\Delta) < \Delta$, utilising $\pi(\Delta)$ processors running in parallel and after a potentially large amount of pre-computation. So, such factors are explicitly incorporated into the puzzle's definitions [10, 39, 23].

Definition 3 (Time-lock Puzzle Security). *A time-lock puzzle is secure if for all λ and Δ, all probabilistic polynomial time adversaries $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ where $\mathcal{A}_1$ runs in total time $O(poly(\Delta, \lambda))$ and $\mathcal{A}_2$ runs in time $\delta(\Delta) < \Delta$ using at most $\pi(\Delta)$ parallel processors, there exists a negligible function $\mu(.)$, such that:*

$$Pr \left[\mathcal{A}_2(pk, \ddot{o}, state) \to b \;\middle|\; \begin{matrix} \text{Setup}(1^\lambda, \Delta) \to (pk, sk) \\ \mathcal{A}_1(1^\lambda, pk, \Delta) \to (s_0, s_1, state) \\ b \xleftarrow{\$} \{0, 1\} \\ \text{GenPuz}(s_b, pk, sk) \to \ddot{o} \end{matrix} \right] \leq \frac{1}{2} + \mu(\lambda)$$

An RSA-based time-lock puzzle construction that realises the above definitions was proposed in [43]. The construction is as follows.

1. **Setup**: TLP.Setup$(1^\lambda, \Delta)$
   (a) Compute $N = q_1 q_2$, where q_i is a large randomly chosen prime number. Then compute Euler's totient function of N, as: $\phi(N) = (q_1 - 1)(q_2 - 1)$
   (b) Set $T = S\Delta$ as the total number of squaring needed to decrypt an encrypted message m, where Δ is the period (in seconds) within which the message should remain private and S is the maximum number of squaring modulo N per second that can be performed by a solver.
   (c) Choose a random key: k for a semantically secure symmetric key encryption that has three algorithms: (GenKey, Enc, Dec)
   (d) Pick a uniformly random value r from $\mathbb{Z}_N^*$
   (e) Compute $a = 2^T \bmod \phi(N)$
   (f) Set $pk = (N, T, r)$ as public key and set $sk = (q_1, q_2, a, k)$ as secret key.
2. **Generate Puzzle**: TLP.GenPuZ(m, pk, sk)
   (a) Encrypt the message using the symmetric key encryption: $o_1 = \text{Enc}(k, m)$
   (b) Encrypt the key: k, as: $o_2 = k + r^a \bmod N$
   (c) Sets: $\ddot{o} = (o_1, o_2)$ as ciphertext or puzzle. Next, output $\ddot{o}$
3. **Solve Puzzle**: TLP.SolvPuz$(pk, \ddot{o})$
   (a) Find b, where $b = r^{2^T} \bmod N$, by using T number of squaring r modulo N

(b) Decrypt the key's ciphertext: $k = o_2 - b \bmod N$

(c) Decrypt the message's ciphertext: $m = \texttt{Dec}(k, o_1)$. Output m

Informally, the time-lock puzzle's security relies on the hardness of factoring problem, the security of the symmetric key encryption, and sequential squaring assumption. We refer readers to Appendix C for more discussion on the construction and its security.

4 Multi-instance Time-lock Puzzle

4.1 Strawman Solution

In the following, we elaborate on the problems that would arise if an existing time-lock puzzle is used directly to handle multiple puzzles at once. Without loss of generality, to illustrate the problems, we use the well-known TLP scheme presented in Section 3.4.

Consider the case where a client wants a server to learn a vector of messages: $\vec{m} = [m_1, ..., m_z]$ at times $[f_1, ..., f_z]$ respectively, where the client is available and online only at an earlier time $f_0 < f_1$. For the sake of simplicity, let $\Delta = f_1 - f_0$ and $\Delta = f_{j+1} - f_j$, where $1 \leq j \leq z$. A naive way to address the problem is that the client uses the TLP to encrypt each message m_j separately, such that it can be decrypted at time f_j if all ciphertexts and public keys are passed on to the server at time t_0. For the server to decrypt the messages on time, it needs to start decrypting *all of them* as soon as the ciphertexts and public keys are given to it.

Parallel Composition Problem. The above naive approach yields two serious issues: (a) imposing a high computation cost, as the server has to perform $S\Delta \sum\limits_{j=1}^{z} j$ squaring to decrypt all messages, and (b) demanding a high level of parallelisation, as each puzzle has to be dealt with separately in parallel to the rest. The issues can be cast as "*parallel composition problem*", where z instances of a puzzle scheme are given at once to a server whose only option, to find solutions on time, is to solve them in parallel[2]. Also, for the client to efficiently compute a_j for each message m_j, where $j > 1$, it has to perform at least one modular multiplication, i.e. $a_j = a_1 a_{j-1} = 2^{jT}$, where $a_1 = 2^T$. In this step, in total $z - 1$ modular multiplications are required to compute all a_j values, for z messages (which is not optimal). Note, we do not see the above issues as previous schemes' flaws, because they were not initially designed for the multi-puzzle setting.

4.2 An Overview of our Solutions

Our key observation is, in the naive approach, the process of decrypting messages has many overlaps leading to a high computation cost. So, by removing the overlaps, we can considerably lower the overall cost both in *puzzle solving* and *puzzle creating* phases. One of our core ideas is to chain the puzzles. While chaining different puzzles may seem a relatively obvious approach to tackle the issues, designing a secure protocol that also can make black-box use of a standard time-lock puzzle scheme, supports public verifiability, and has low costs is challenging. In our solution, a client first encrypts the message that is supposed to be decrypted after the rest and embeds the information needed for decrypting it into the ciphertext of the message that will be decrypted before

---

[2] It should not be confused with the "universally composable" notion put forth in [14].

that message. In other words, the client integrates the information (i.e. a part of public keys) needed to decrypt message m_j into the ciphertext related to message m_{j-1}. In this case, the server after learning message m_{j-1} at time f_{j-1} learns the public key needed to perform the sequential squaring to decrypt the next message: m_j. This means after fully decrypting m_{j-1}, the server starts squaring sequentially to decrypt m_j

Addressing Parallel Composition Problem. The above approach solves the parallel composition problem for two main reasons. First, the total number of squaring required to decrypt all z messages is now much lower, i.e. $S\Delta z$, and is equivalent to the number of squaring needed to solve only the last puzzle, i.e. z-th one. Second, it does not call for high parallelisation. Because now the server does not need to deal with all of the puzzles in parallel; instead, it solves them sequentially one after another.

Adding Efficient Publicly Verifiable Algorithm. To let the scheme support efficient public verifiability, we use the following novel trick. The client uses a commitment scheme to commit to every message: m_i and publishes the commitment. Then, it uses the time-lock encryption to encrypt the commitment's opening, i.e. a combination of m_i and a random value. But, unlike the traditional commitment, the client does not open the commitment itself. Instead, the server does that, after it discovers the puzzle's solution. When it finds a solution, it decodes the solution to find the opening and sends it to the public who can check the solution correctness. So, to verify a solution's correctness, a verifier only needs to run the commitment's verification algorithm that is: (a) publicly verifiable, and (b) efficient. It can be built in the random oracle or the standard model.

The approach also allows the client at the setup to compute only a single $a = 2^T$ reusable for all z puzzles, imposing only $O(1)$ cost.

4.3 Multi-instance Time-lock Puzzle Definition

In this section, we provide a formal definition of a multi-instance time-lock puzzle. Our starting point is the time-lock puzzle definition, i.e. Definition 2, but we extend it from several perspectives, so it can: (a) handle multiple solutions/messages in setup, (b) produce multiple puzzles for the messages, (c) solve the puzzles given the puzzles and public parameters, and (d) support public verifiability. In the following, we provide the formal definition of a multi-instance time-lock puzzle.

Definition 4 (Multi-instance Time-lock Puzzle). *A multi-instance time-lock puzzle has the following five algorithms and satisfies completeness and efficiency properties.*

- ***Algorithms:***
  - $\texttt{Setup}(1^\lambda, \Delta, z) \rightarrow (pk, sk, \vec{d})$: *a probabilistic algorithm that takes as input security: 1^λ and time: Δ parameters and the total number of solutions/puzzles: z. Let $j\Delta$ be a time period after which j-th solution is found. It outputs public-private key pair: (pk, sk) and a vector of fixed size secret witnesses: $\vec{d}$*
  - $\texttt{GenPuz}(\vec{m}, pk, sk, \vec{d}) \rightarrow \ddot{o}$: *a probabilistic algorithm that takes as an input a message vector: $\vec{m} = [m_1, ..., m_z]$, the public-private key pair: (pk, sk), and the witness vector: $\vec{d}$. It outputs $\ddot{o} : (\vec{o}, \vec{h})$, where $\vec{o}$ is a puzzle vector, and $\vec{h}$ is a commitment vector. Each j-th element in vectors $\vec{o}$ and $\vec{h}$ corresponds to a solution s_j of the form: $s_j = m_j || d_j$*

- $\texttt{SolvPuz}(pk, \overrightarrow{\boldsymbol{o}}) \to \overrightarrow{\boldsymbol{s}}$: *a deterministic algorithm that takes as input the public key: pk and puzzle vector: $\overrightarrow{\boldsymbol{o}}$. It outputs a solution vector: $\overrightarrow{\boldsymbol{s}}$*
- $\texttt{Prove}(pk, s_j) \to \ddot{p}_j$: *a deterministic algorithm that takes the public key: pk and a solution: $s_j \in \overrightarrow{\boldsymbol{s}}$. It outputs a proof, $\ddot{p}_j : (m_j, d_j)$*
- $\texttt{Verify}(pk, \ddot{p}_j, h_j) \to \{0, 1\}$: *a deterministic algorithm that takes public key: pk, proof: $\ddot{p}_j$ and commitment: $h_j \in \overrightarrow{\boldsymbol{h}}$. It outputs 0 if it rejects, or 1 if it accepts.*
  - **Completeness**: *for any honest prover and verifier, it always holds that:*
    - $\texttt{SolvPuz}(pk, [o_1, ..., o_j]) = [s_1, ..., s_j]$, *for every j, $1 \leq j \leq z$*
    - $\texttt{Verify}(pk, \texttt{Prove}(pk, s_j), h_j) \to 1$
  - **Efficiency**: *the run-time of algorithm $\texttt{SolvPuz}(pk, [o_1, ..., o_j]) = [s_1, ...s_j]$ is bounded by: $poly(j\Delta, \lambda)$, where $poly(.)$ is a fixed polynomial and $1 \leq j \leq z$*

Informally, a multi-instance time-lock puzzle is secure if it satisfies two properties: a solution's *privacy* and *validity*. The former requires its j-th solution to remain hidden from all adversaries running in parallel within time period: $j\Delta$, while the latter one requires that it is infeasible for a PPT adversary to come up with an invalid solution and passes the verification. The two properties are formally defined in Definitions 5 and 6.

Definition 5 (Multi-instance Time-lock Puzzle's Solution-Privacy). *A multi-instance time-lock puzzle is privacy-preserving if for all λ and Δ, any number of puzzle: $z \geq 1$, any pair of randomised algorithm $\mathcal{A} : (\mathcal{A}_1, \mathcal{A}_2)$, where $\mathcal{A}_1$ runs in time $O(poly(j\Delta, \lambda))$ and $\mathcal{A}_2$ runs in time $\delta(j\Delta) < j\Delta$ using at most $\pi(\Delta)$ parallel processors, there exists a negligible function $\mu(.)$, such that:*

$$Pr\begin{bmatrix} \mathcal{A}_2(pk, \ddot{o}, state) \to \ddot{a} \\ s.t. \\ \ddot{a} : (b_i, i) \\ m_{b_i, i} = m_{b_j, j} \end{bmatrix} \begin{matrix} \texttt{Setup}(1^\lambda, \Delta, z) \to (pk, sk, \overrightarrow{\boldsymbol{d}}) \\ \mathcal{A}_1(1^\lambda, pk, z) \to (\overrightarrow{\boldsymbol{m}}, state) \\ [b_1, ..., b_z], b_j \xleftarrow{\$} \{0, 1\} \\ \texttt{GenPuz}(\overrightarrow{\boldsymbol{m}}', pk, sk, \overrightarrow{\boldsymbol{d}}) \to \ddot{o} \end{matrix} \Bigg] \leq \frac{1}{2} + \mu(\lambda)$$

where $\overrightarrow{\boldsymbol{m}} : [(m_{0,1}, m_{1,1}), ..., (m_{0,z}, m_{1,z})]$, $\overrightarrow{\boldsymbol{m}}' : (m_{b_1,1}, ..., m_{b_z,z})$, $1 \leq j \leq z$ and $1 \leq i \leq z$

The definition above also ensures the solutions to appear after j-th one, remain hidden from the adversary with a high probability, as well. Similar to [10, 39, 23], it captures that even if $\mathcal{A}_1$ computes on the public parameters for a polynomial time, $\mathcal{A}_2$ cannot find j-th solution in time $\delta(j\Delta) < j\Delta$ utilising $\pi(\Delta)$ parallel processors, with a probability significantly greater than $\frac{1}{2}$. As highlighted in [10], we can set $\delta(\Delta) = (1 - \epsilon)\Delta$ for a small ϵ, where $0 < \epsilon < 1$

Definition 6 (Multi-instance Time-lock Puzzle's Solution-Validity). *A multi-instance time-lock puzzle preserves a solution validity, if for all λ and Δ, any number of puzzles: $z \geq 1$, all probabilistic polynomial-time adversaries $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$ that run in time $O(poly(\Delta, \lambda))$ there is negligible function $\mu(.)$, such that:*

$$Pr\begin{bmatrix} \mathcal{A}_2(pk, \overrightarrow{\boldsymbol{s}}, \ddot{o}, state) \to a \\ s.t. \\ a : (j, \ddot{p}_j, \ddot{p}') \\ \ddot{p}_j : (m_j, d_j), \ddot{p}' : (m', d') \\ m_j \in \overrightarrow{\boldsymbol{m}}, d_j \in \overrightarrow{\boldsymbol{d}}, m \neq m' \\ \texttt{Verify}(pk, \ddot{p}, h_j) = 1 \\ \texttt{Verify}(pk, \ddot{p}', h_j) = 1 \end{bmatrix} \begin{matrix} \texttt{Setup}(1^\lambda, \Delta, z) \to (pk, sk, \overrightarrow{\boldsymbol{d}}) \\ \mathcal{A}_1(1^\lambda, pk, \Delta, z) \to (\overrightarrow{\boldsymbol{m}}, state) \\ \texttt{GenPuz}(\overrightarrow{\boldsymbol{m}}, pk, sk, \overrightarrow{\boldsymbol{d}}) \to \ddot{o} \\ \texttt{SolvPuz}(pk, \overrightarrow{\boldsymbol{o}}) \to \overrightarrow{\boldsymbol{s}} \end{matrix} \Bigg] \leq \mu(\lambda)$$

where $\overrightarrow{\boldsymbol{m}} = [m_1, ..., m_z]$, and $h_j \in \overrightarrow{\boldsymbol{h}} \in \ddot{o}$

Definition 7 (Multi-instance Time-lock Puzzle Security). *A multi-instance time-lock puzzle scheme is secure if it meets solution-privacy and solution-validity properties.*

4.4 Chained Time-lock Puzzle (C-TLP) Protocol

In this section, we present the chained time-lock puzzle (C-TLP), an instantiation of the multi-instance time lock puzzle. Since we have already presented an outline of C-TLP (in Section 4.2), in this section we present C-TLP protocol in detail. Recall, a client wants a server to learn a vector of messages: $\vec{m} = [m_1, ..., m_z]$ at times $[f_1, ..., f_z]$ respectively, where the client is available and online only at an earlier time $f_0 < f_1$. Also, the client wants to ensure that anyone can validate a solution found by the server, i.e. supports public verifiability. For the sake of simplicity, let $\Delta = f_1 - f_0$ and $\Delta = f_{j+1} - f_j$, where $1 \leq j \leq z$ and $T = S\Delta$. Below, we provide C-TLP protocol. We refer readers to Appendix D for further remarks on the protocol.

1. **Setup**: $\texttt{Setup}(1^\lambda, \Delta, z)$.
   (a) Call: $\texttt{TLP.Setup}(1^\lambda, \Delta) \to (\hat{pk}, \hat{sk})$, s.t. $\hat{pk} = (N, T, r_1)$ and $\hat{sk} = (q_1, q_2, a, k_1)$
   (b) Pick $z - 1$ fixed size random generators: $\vec{r} = [r_2, ..., r_z]$ from $\mathbb{Z}_N^*$
   (c) Pick $z - 1$ random keys: $[k_2, ..., k_z]$ for a symmetric key encryption. Let $\vec{k} = [k_1, ..., k_z]$, where $k_1 \in \hat{sk}$. Also, pick z fixed size sufficiently large random values: $\vec{d} = [d_1, ..., d_z]$, e.g. $|d_j| = 128$-bit or 1024-bit depending on the choice of a commitment scheme.
   (d) Set $pk = (\text{aux}, N, T, r_1)$ as public key. Set $sk = (q_1, q_2, a, \vec{k}, \vec{r}, \vec{d})$ as secret key. Note, aux contains a cryptographic hash function's description and the size of the random values. Also, note that all generators, except r_1 are kept secret. Output pk and sk
2. **Generate Puzzle**: $\texttt{GenPuz}(\vec{m}, pk, sk)$
   Encrypt the messages, starting with $j = z$, in descending order. $\forall j, z \geq j \geq 1$:
   (a) Set $pk_j = (N, T, r_j)$ and $sk_j = (q_1, q_2, a, k_j)$. Note, if $j = 1$ then $r_j \in pk$; otherwise (when $j > 1$), $r_j \in \vec{r}$
   (b) Generate a puzzle (or ciphertext pair):
      - if $j = z$, then run: $\texttt{TLP.GenPuZ}(m_j || d_j, pk_j, sk_j) \to \ddot{o}_j$
      - otherwise, run: $\texttt{TLP.GenPuZ}(m_j || d_j || r_{j+1}, pk_j, sk_j) \to \ddot{o}_j$
   (c) Commit to each message, e.g. $\texttt{H}(m_j || d_j) = h_j$ and output: h_j
   (d) Output: $\ddot{o}_j = (o_{j,1}, o_{j,2})$ as puzzle (or ciphertext pair).
   By the end of this phase, vectors of puzzles: $\vec{o} = [\ddot{o}_1, ..., \ddot{o}_z]$ and commitments: $\vec{h} = [h_1, ... h_z]$ are generated. All public parameters and puzzles are given to a server at time $t_0 < t_1$, where $\Delta = f_1 - f_0$
3. **Solve Puzzle**: $\texttt{SolvPuz}(pk, \vec{o})$
   Decrypt the messages, starting with $j = 1$, in ascending order. $\forall j, 1 \leq j \leq z$:
   (a) If $j = 1$, then set $r_j = r_1$, where $r_1 \in pk$; Otherwise, set $r_j = u$
   (b) Set $pk_j = (N, T, r_j)$
   (c) Run: $\texttt{TLP.SolvPuz}(pk_j, \ddot{o}_j) \to x_j$, where $\ddot{o}_j \in \vec{o}$
   (d) Parse x_j. Note that if $j < z$ then $x_j = m_j || d_j || r_{j+1}$; otherwise, we have $x_j = m_j || d_j$. Therefore, x_j is parsed as follows.

- if $j < z$:
    - i. Parse $m_j||d_j||r_{j+1}$ into $m_j||d_j$ and $u = r_{j+1}$
    - ii. Output $s_j = m_j||d_j$
- otherwise (when $j = z$), output $s_j = x_j = m_j||d_j$

4. **Prove**: $\text{Prove}(pk, s_j)$. Parse s_j into $\ddot{p}_j : (m_j, d_j)$, and send the pair to the verifier.
5. **Verify**: $\text{Verify}(pk, \ddot{p}_j, h_j)$. Verifies the commitment, $\text{H}(m_j, d_j) \overset{?}{=} h_j$. If passed, accept the solution and output 1; otherwise, reject it and output 0

Theorem 1 (C-TLP Security). *C-TLP is a secure multi-instance time-lock puzzle.*

Proof (Outline). The proof of Theorem 1 relies on the security of the TLP, symmetric key encryption, and commitment schemes. It is also based on the fact that the probability to find a certain random generator is negligible. It shows both C-TLP's solution privacy (due to security of the above three schemes) and validity (due to the security of the commitment) are satisfied. We refer readers to Appendix E for detailed proof.

4.5 Cost Analysis Table

We summarize C-TLP's cost analysis in Table 1. It considers a generic setting where the protocol deals with z puzzles. We refer readers to Appendix F for detailed analysis.

Table 1: C-TLP's Detailed Cost Breakdown

(a) Computation Cost

Protocol	Operation	Protocol Function			Complexity
		GenPuz	SolvPuz	Verify	
C-TLP	Exp.	$z+1$	Tz	—	$O(Tz)$
	Add. or Mul.	z	z	—	
	Commitment	z	—	z	
	Sym. Enc	z	z	—	

(b) Communication Cost (in bit)

Protocol	Model	Client	Server	Complexity
C-TLP	Standard	$3200z$	$1524z$	$O(z)$
	R.O.	$2432z$	$628z$	

5 Smarter Outsourced PoR (SO-PoR) Utilising C-TLP

As discussed in Section 2, the existing outsourced PoR's have serious shortcomings, e.g. having high costs, not supporting real-time detection, or suffering from the lack of a fair payment mechanism. In this section, we present SO-PoR to addresses them.

5.1 SO-PoR Overview

SO-PoR uses a unique combination of (a) homomorphic MAC-based PoR [46], (b) C-TLP, (c) a smart contract, (d) a pre-computation technique, and (e) blockchain-based random extraction beacon [1, 3]. It uses the MAC-based PoR, due to its high efficiency. Since the MAC's are privately verifiable and secret verification keys are needed to check PoR proofs, SO-PoR also uses C-TLP to efficiently make them publicly verifiable. In this case, C-TLP encapsulates the verification keys and reveals each of them to verifiers only after a certain time. SO-PoR also utilises a smart contract which acts as a public verifier on the client's behalf to verify proofs and pay an honest server. The

pre-computation technique allows the client at setup to generate a constant number of *disposable* homomorphic MAC's for each verification. The combination of disposable homomorphic MAC's and C-TLP makes it possible to (a) use a smart contract and (b) take advantage of MAC's efficiency in the setting where public verifiability is needed. This combination has applications beyond PoR. A blockchain-based random extraction beacon allows the server to independently derive a set of unpredictable random values from the blockchain such that the values' correctness is publicly verifiable.

At a high-level SO-PoR works as follows. The client encodes its file using an error-correcting code and for each j-th verification it does the following. It picks two random keys: (v_j, l_j) of a PRF. It uses v_j to generate c random blocks' indices, i.e. challenged blocks. It utilises l_j to generate a disposable MAC on each challenged block. It also uses C-TLP to make two puzzles, one that encapsulates v_j, and another that encapsulates l_j. It deposits enough coins to cover z successful PoR verifications in a smart contract. The client sends the encoded file, tags and the puzzles to the server. When j-th PoR proof is needed, the server manages to discover key v_j that lets it determine which file blocks are challenges. The server also uses the beacon to extract a set of random values from the blockchain. Using the MAC's, challenged blocks, and beacon's outputs, the server generates a compact PoR proof. The server sends the proof to the contract. After that, it can delete the related disposable MAC's. For the same verification, after a fixed time, it manages to find the related MAC's verification key: l_j. It sends the key to the server who checks the correctness of l_j and PoR proof. If the contract accepts all proofs, then it pays the server for j-th verification; otherwise, it notifies the client.

5.2 SO-PoR Model Overview

SO-PoR model is built upon the traditional PoR paradigm [46] which is a challenge-response protocol where a server proves to an honest client that its file is retrievable (see Appendix G for a formal definition of the PoR). In SO-PoR, however, a client may not be available for verification. So, it wants to delegate a set of verifications that it cannot carry out. Informally, in this setting, it (in addition to file retrievability) must have three guarantees: (a) *verification correctness*: every verification is performed honestly, so the client can trust the verification's result without redoing it, (b) *real-time detection*: the client is notified in almost real-time when a proof is rejected, and (c) *fair payment*: in every verification, the server is paid only if a proof is accepted. In SO-PoR, three parties are involved: an honest client, potentially malicious server and a standard smart contract. SO-PoR also, analogous to [46], allows a client to perform the verification itself, when it is available. We present our formal definition of SO-PoR in Appendix H.

5.3 SO-PoR Protocol

This section presents SO-PoR protocol in detail, followed by the rationale behind it.

1. ***Client-side Setup***.
   (a) ***Gen. Public and Private Keys***: Picks a fresh key: $\hat{k}$ and two vectors of keys: $\vec{v}$ and $\vec{l}$, where each vector contains z fresh keys. It picks a large prime number: p whose size is determined by a security parameter, i.e. $|p| = \iota$. Moreover, it runs $\texttt{Setup}(.)$ in C-TLP scheme to generate a key pair: (pk, sk)

(b) ***Gen. Other Public Parameters***: Sets c to the total number of blocks challenged in each verification. It defines parameters: w and g, where w is an index of a future block: $\mathcal{B}_w$ in a blockchain that will be added to the blockchain (permanent state) at about the time first delegated verification will be done, and g is a security parameter referring to the number of blocks (in a row) starting from w. It sets z: the total number of verifications, $||\boldsymbol{F}||$: file bit size, Δ_1: the maximum time is taken by the server to generate a proof, Δ_2: time window in which a message is (sent by the server and) received by the contract, and e amount of coins paid to the server for each successful verification. Sets $\hat{pk} : (pk, e, g, w, p, c, z, \Delta_1, \Delta_2)$

(c) ***Sign and Deploy Smart Contract***: Signs and deploys a smart contract: $\mathcal{SC}$ to a blockchain. It stores public parameters: $(z, ||\boldsymbol{F}||, \Delta_1, \Delta_2, c, g, p, w)$, on the contract. It deposits ez coins to the contract. Then, it asks the server to sign the contract. The server signs if it agrees on all parameters.

2. ***Client-side Store***.

(a) ***Encode File***: Splits an error-corrected file, e.g. under Reed-Solomon codes, into n blocks; $\boldsymbol{F} : [F_1, ..., F_n]$, where $F_i \in \mathbb{F}_p$

(b) ***Gen. Permanent Tags***: Using the key: $\hat{k}$, it computes n pseudorandom values: r_i and single value: α, as follows.
$$\alpha = \text{PRF}(\hat{k}, n+1) \bmod p, \qquad \forall i, 1 \le i \le n : r_i = \text{PRF}(\hat{k}, i) \bmod p$$
It uses the pseudorandom values to compute tags for the file blocks.
$$\forall i, 1 \le i \le n : \sigma_i = r_i + \alpha \cdot F_i \bmod p$$
So, at the end of this step, a set of tags are generated, $\sigma : \{\sigma_1, ..., \sigma_n\}$

(c) ***Gen. Disposable Tags***: For j-th verification ($1 \le j \le z$):

i. chooses the related key: $v_j \in \vec{v}$ and computes c pseudorandom indices.
$$\forall b, 1 \le b \le c : x_{b,j} = \text{PRF}(v_j, b) \bmod n$$

ii. picks the corresponding key: $l_j \in \vec{l}$ and computes c pseudorandom values: $r_{b,j}$ and single value: α_j, as follows.
$$\alpha_j = \text{PRF}(l_j, c+1) \bmod p, \qquad \forall b, 1 \le b \le c : r_{b,j} = \text{PRF}(l_j, b) \bmod p$$

iii. generates c disposable tags.
$$\forall b, 1 \le b \le c : \sigma_{b,j} = r_{b,j} + \alpha_j \cdot F_y \bmod p$$
where $y = x_{b,j}$. At the end of this step, a set σ_j of c tags are computed, $\sigma_j : \{\sigma_{1,j}, ..., \sigma_{c,j}\}$

(d) ***Gen. Puzzles***: Sets $\overrightarrow{m} = [v_1, l_1, ..., v_z, l_z]$ and then encrypts the vector's elements, by running: $\text{GenPuz}(\overrightarrow{m}, pk, sk)$ in C-TLP scheme. This yields a puzzle vector: $[(V_1, L_1), ..., (V_z, L_z)]$ and a commitment vector: $\overrightarrow{h}$. The encryption is done in such a way that in each j-th pair, V_j will be fully decrypted at times t_j and L_j will be decrypted at time t'_j, where $t_j + \Delta_1 + \Delta_2 \le t'_j < t_{j+1}$

(e) ***Outsource File***: Stores $\boldsymbol{F}, n, \hat{pk}, \{\sigma, \sigma_1, ..., \sigma_z, (V_1, L_1), ..., (V_z, L_z)\}$ on the server. Also, it stores $\overrightarrow{h}$ on the smart contract.

3. ***Cloud-Side Proof Generation***. For j-th verification ($1 \le j \le z$), the cloud:

(a) ***Solve Puzzle and Regen. Indices***: Receives and parses the output of $\text{SolvPuz}(.)$ in C-TLP, to extract v_j, at time t_j. Using v_j, it regenerates c pseudorandom indices.
$$\forall b, 1 \le b \le c : x_{b,j} = \text{PRF}(v_j, b) \bmod n$$

(b) ***Extract Key***: Extracts a seed: u_j, from the blockchain as follows: $u_j = \text{H}(\mathcal{B}_\gamma || ... || \mathcal{B}_\varsigma)$, where $\gamma = w + (j-1) \cdot g$ and $\varsigma = w + j \cdot g$

(c) *Gen. PoR*: Generates a PoR proof.

$$\mu_j = \sum_{b=1}^{c} \mathtt{PRF}(u_j, b) \cdot F_y \bmod p, \quad \xi_j = \sum_{b=1}^{c} \mathtt{PRF}(u_j, b) \cdot \sigma_{b,j} \bmod p$$

where y is a pseudorandom index: $y = x_{b,j}$

(d) *Register Proofs*: Sends the PoR proof: (μ_j, ξ_j) to the smart contract within Δ_1

(e) *Solve Puzzle and Regen. Verification Key*: Receives and parses the output of algorithm $\mathtt{SolvPuz}(.)$ in C-TLP to extract l_j, at time t'_j. Also, it runs $\mathtt{Prove}(.)$ in C-TLP, to generate a proof: $\ddot{p}_j$, of l_j's correctness. It sends $\ddot{p}_j$ (containing l_j) to the contract, so it can be received by the contract within Δ_2

4. *Smart Contract-Side Verification*. For j-th verification ($1 \le j \le z$), the contract:

(a) *Check Arrival Time*: checks the arrival time of the decrypted values sent by the server. In particular, it checks, if (μ_j, ξ_j) was received in the time window: $(t_j, t_j + \Delta_1 + \Delta_2]$ and whether l_j was received in the time window: $(t'_j, t'_j + \Delta_2]$

(b) *Verify Puzzle Solution*: runs $\mathtt{Verify}(.)$ in C-TLP to verify $\ddot{p}_j$ (i.e. checks the correctness of $l_j \in \ddot{p}_j$). If approved, then it regenerates the seed: $u_j = \mathtt{H}(\mathcal{B}_\gamma || ... || \mathcal{B}_\zeta)$, where $\gamma = w + (j-1) \cdot g$ and $\zeta = w + j \cdot g$

(c) *Verify PoR*: regenerates the pseudorandom values and verifies the PoR proof.

$$\xi_j \overset{?}{=} \mu_j \cdot \mathtt{PRF}(l_j, c+1) + \sum_{b=1}^{c} (\mathtt{PRF}(u_j, b) \cdot \mathtt{PRF}(l_j, b)) \bmod p \qquad (1)$$

(d) *Pay*: if Equation 1 holds, pays and asks the server to delete all disposable tags for this verification, i.e. σ_j

If either check fails, it aborts and notifies the client.

5. *Client-server PoR*: When the client is online, it can interact with the server to check its data availability too. In particular, it sends c random challenges and random indices to the server who computes POR using only: (a) the messages sent by the client in this step, (b) the file: F, and (c) the tags: $\sigma_i \in \sigma$, generated in step 2b. The proof generation and verification are similar to the MAC-based schemes, e.g. [46].

Theorem 2. *SO-PoR protocol is secure if the MAC's are unforgeable,* $\mathtt{PRF}(.)$ *is a secure pseudorandom function, the blockchain and C-TLP protocol are secure, and* $\mathtt{H}(\mathcal{B}_\gamma || ... || \mathcal{B}_\zeta)$ *outputs an unpredictable random value (where* $\zeta - \gamma$ *is a security parameter).*

Remark 1. We refer readers to Appendix I for the above theorem's proof.

Remark 2. In SO-PoR, for a security reason the server must record j-th PoR proof in the contract before l_j is recovered. Also, the way disposable tags are generated in SO-PoR differs from those computed in previous PoR schemes, despite having similarities structure-wise. Moreover, with slight adjustments, we can reduce the contract-side storage cost to constant. For more details, we refer readers to Appendix J which also explains why strawman approaches are not suitable substitutes for SO-PoR.

5.4 Evaluation

In this section, we provide a summary of comparisons between SO-PoR and outsourced PoRs [3, 53, 5]. Among the two protocols proposed in [5] we only consider "basic PoSt" as it supports public verifiability. Briefly, in terms of property, only SO-PoR offers

an explicit solution for real-time detection and fair payment. In terms of computation cost, the verification algorithm in SO-PoR is much faster than the other three protocols; Specifically, when $c = 460$, SO-PoR verification[3] requires about 4.5 times fewer computation than the verification required in the fastest outsourced PoR [3]. Also, [3] has the worst store cost, which is much higher than that of SO-PoR; e.g. for a 1-GB file, SO-PoR requires over 46×10^5 times fewer exponentiations than [3] needs in the same phase. SO-PoR and [5] require a server to solve puzzles but the other two protocols do not need that. Also, the I/O cost and proof complexity of all protocols are $O(1)$ except [5] whose I/O cost and proof complexity are $O(\log n)$. The server-side bandwidth of SO-PoR is much lower than the rest; for instance, for 1-GB file and $z = 100$ verifications, a server in SO-PoR requires 9×10^4, 7 and 1729 times fewer bits than those required in [3], [53] and [5] respectively. A client in SO-PoR has a higher bandwidth than the rest of the protocols (but this cost is one-off). Thus, SO-PoR offers additional properties, it has lower verification cost and lower server-side bandwidth than the rest while its other costs remain reasonable. Table 2 outlines the cost comparison results. For a full analysis, we refer readers to Appendix K where we also compare SO-PoR costs with the costs of the most efficient traditional PoR [46].

Table 2: Outsourced PoR's Cost Comparison. In the table, z is the total number of verifications, c is the number of challenges for each verification, n is the total number of file blocks, $c' = (0.1)c$, and $||\boldsymbol{F}||$ is a file bit size.

Protocols	Operation	Computation Cost				Communication Cost			
		Store	SolvPuz	Prove	Verify	Client	Server	Verifier	Proof Size
SO-PoR	Exp.	$z+1$	Tz	—	—	$128(n+$ $cz + 19z)$	$884z$	—	$O(1)$
	Add. or Mul.	$2(n+cz)$	z	$4cz$	$2z(1+c)$				
[3]	Exp.	$9n$	—	—	—	$128n$	$256z+$ $||\boldsymbol{F}||$	$4672n+$ $256z$	$O(1)$
	Add. or Mul.	$10n$	—	$4z(c+c')$	$z(9c+3)$				
[53]	Exp.	—	—	$z(3+c)$	$6z$	$2048n$	$6144z$	—	$O(1)$
	Add. or Mul.	$4n$	—	$2z(3c+4)$	$2cz$				
	Pairing	—	—	$7z$	—				
[5]	Exp.	—	$3Tz$	—	$3z$	128	$128cz \log n+$ $4096z$	—	$O(\log n)$
	Add. or Mul.	—	Tz	—	—				

6 C-TLP as Efficient Variant of VDF

There are cases where a client wants a server to learn distinct random challenges, at different points in time within a certain period, without the client's involvement in that period. Such challenges can let the server generate certain proofs that include but are not limited to the *continuous availability of services*, such as data storage or secure hardware. The obvious candidates that can meet the above needs are VDF (if public verifiability is desirable) and TDF (if private verifiability suffices). PoSt protocols in [5] are two examples of the above cases. We observed that in these cases, VDF/TDF can be replaced with C-TLP to gain better efficiency. The idea is that the client computes random challenges, encodes them into C-TLP puzzles and sends them to the server who can eventually solve each puzzle, extract a subset of challenges and use them for the related proof scheme while letting the public efficiently verify the solutions' correctness.

---

[3] As shown in [4], to ensure 99% of file blocks is retrievable, it suffices to set $c = 460$.

To illustrate the efficiency gain, we compare C-TLP performance with the two current VDF functions [51, 10]. Table 3 summarises the result. The cost analysis considers the generic setting where z outputs are generated. Among several VDF schemes proposed in [10], we focus on the one that uses sequential squaring, as it is more efficient than the other schemes in [10]. As the table indicates, the overall cost of [10] in each of the three phases is much higher than C-TLP and [51]. Now, we compare the computation cost of C-TLP with [51]. At setup, a client in C-TLP performs at most $3z + 1$ more exponentiations than it does in [51]. But, at both prove and verify phases, C-TLP outperforms [10], especially when they are in the same model. In particular, at the prove phase, C-TLP, in both models, requires Tz fewer multiplications than [51] does. Also, in the same phase, it requires 3 times fewer exponentiations than [51]. In the verify phase, when C-TLP is in the standard model, it has a slightly lower cost than [51] has in the random oracle model. However, when both of them are in the random oracle model, C-TLP has a much lower cost, as it requires no exponentiations whereas [51] needs $3z$ exponentiations. Hence, C-TLP supports both standard and random oracle models and in both paradigms, it outperforms the fasted VDF, i.e. [51], designed in the random oracle model. Furthermore, the proof size in C-TLP is 3.2 and 6.5 times shorter than [10] and [51] respectively, when they are in the same model. In Appendix L, we also show how C-TLP can be employed in the PoSt protocols [5] to reduce costs.

Table 3: VDF's Cost Comparison

Protocols	Model	Operation	Computation Cost			Proof size (bit)
			Setup	Prove	Verify	
C-TLP	Standard	Exp.	$3z + 1$	Tz	$2z$	$1524z$
		Mul.	z	$-$	z	
	R.O.	Exp.	$z + 1$	Tz	$-$	$628z$
		Mul.	$-$	$-$	$-$	
[10]	R.O.	Exp.	2^{30}	Tz	z	$2048z$
		Mul.	$-$	$2z \cdot 2^{30}$	$2z \cdot 2^{30}$	
[51]	R.O.	Exp.	$-$	$3Tz$	$3z$	$4096z$
		Mul.	$-$	Tz	$-$	

7 Conclusion

Time-lock puzzles are important cryptographic protocols with various applications. Nevertheless, existing puzzle schemes are not suitable to deal with multiple puzzles at once. In this work, we put forth the concept of composing multiple puzzles, where given puzzles composition at once, a server can find one puzzle's solution after another. This process does not require the server to deal with all of them in parallel which reliefs the server from having numerous parallel processors and allows it to save considerable computation overhead. We proposed a candidate construction: chained time-lock puzzle (C-TLP) that possesses the aforementioned features. Furthermore, C-TLP is equipped with an efficient verification algorithm publicly executable. We also illustrated how to use C-TLP to construct an efficient outsourced proofs of retrievability scheme that supports *real-time detection* and *fair payment* while keeping its costs considerably lower than the state of the art protocols. Moreover, we showed how VDF's in certain settings can be replaced with C-TLP to gain considerable cost improvement.

References

1. Abadi, A., Ciampi, M., Kiayias, A., Zikas, V.: Timed signatures and zero-knowledge proofs -timestamping in the blockchain era-. IACR Cryptology ePrint Archive 2019, 644 (2019)
2. Armknecht, F., Barman, L., Bohli, J., Karame, G.O.: Mirror: Enabling proofs of data replication and retrievability in the cloud. In: 25th USENIX Security 16
3. Armknecht, F., Bohli, J.M., Karame, G.O., Liu, Z., Reuter, C.A.: Outsourced proofs of retrievability. In: CCS'14
4. Ateniese, G., Burns, R.C., Curtmola, R., Herring, J., Kissner, L., Peterson, Z.N.J., Song, D.X.: Provable data possession at untrusted stores. In: CCS'07
5. Ateniese, G., Chen, L., Etemad, M., Tang, Q.: Proof of storage-time: Efficiently checking continuous data availability. In: NDSS'20
6. Ateniese, G., Pietro, R.D., Mancini, L.V., Tsudik, G.: Scalable and efficient provable data possession. In: SECURECOMM'08
7. Badertscher, C., Maurer, U., Tschudi, D., Zikas, V.: Bitcoin as a transaction ledger: A composable treatment. In: Katz, J., Shacham, H. (eds.) CRYPTO'17
8. Banerjee, P., Nikam, N., Ruj, S.: Blockchain enabled privacy preserving data audit. CoRR abs/1904.12362 (2019)
9. Bellare, M., Canetti, R., Krawczyk, H.: Keying hash functions for message authentication. In: Advances in Cryptology - CRYPTO '96
10. Boneh, D., Bonneau, J., Bünz, B., Fisch, B.: Verifiable delay functions. In: Shacham, H., Boldyreva, A. (eds.) CRYPTO'18
11. Boneh, D., Naor, M.: Timed commitments. In: Bellare, M. (ed.) CRYPTO 2000
12. Brakerski, Z., Döttling, N., Garg, S., Malavolta, G.: Leveraging linear decryption: Rate-1 fully-homomorphic encryption and time-lock puzzles. In: TCC'19,
13. Campanelli, M., Gennaro, R., Goldfeder, S., Nizzardo, L.: Zero-knowledge contingent payments revisited: Attacks and payments for services. In: CCS'17
14. Canetti, R.: Universally composable security: A new paradigm for cryptographic protocols. In: 42nd Annual Symposium on Foundations of Computer Science, FOCS 2001, 14-17 October 2001, Las Vegas, Nevada, USA. pp. 136–145. IEEE Computer Society (2001)
15. Cash, D., Küpçü, A., Wichs, D.: Dynamic proofs of retrievability via oblivious ram. J. Cryptol. (2017)
16. Chen, H., Deviani, R.: A secure e-voting system based on RSA time-lock puzzle mechanism. In: BWCCA'12 ,
17. Dwork, C., Naor, M.: Pricing via processing or combatting junk mail. In: Brickell, E.F. (ed.) CRYPTO'92
18. Erway, C.C., Küpçü, A., Papamanthou, C., Tamassia, R.: Dynamic provable data possession. In: CCS'09
19. Etemad, M., Küpçü, A.: Transparent, distributed, and replicated dynamic provable data possession. In: Jacobson, M., Locasto, M., Mohassel, P., Safavi-Naini, R. (eds.) ACNS'13
20. Francati, D., Ateniese, G., Faye, A., Milazzo, A.M., Perillo, A.M., Schiatti, L., Giordano, G.: Audita: A blockchain-based auditing framework for off-chain storage. CoRR'19
21. Garay, J.A., Jakobsson, M.: Timed release of standard digital signatures. In: Blaze, M. (ed.) FC'02
22. Garay, J.A., Kiayias, A., Leonardos, N.: The bitcoin backbone protocol with chains of variable difficulty. In: Katz, J., Shacham, H. (eds.) CRYPTO'17
23. Garay, J.A., Kiayias, A., Panagiotakos, G.: Iterated search problems and blockchain security under falsifiable assumptions. IACR Cryptology ePrint Archive (2019)
24. Groza, B., Petrica, D.: On chained cryptographic puzzles. In: SACI'06,

25. Hao, K., Xin, J., Wang, Z., Jiang, Z., Wang, G.: Decentralized data integrity verification model in untrusted environment. In: APWeb-WAIM'18
26. Juels, A., Jr., B.S.K.: Pors: proofs of retrievability for large files. In: CCS'07
27. Kamara, S.: Proofs of storage: Theory, constructions and applications. In: CAI'13
28. Karame, G., Capkun, S.: Low-cost client puzzles based on modular exponentiation. In: Gritzalis, D., Preneel, B., Theoharidou, M. (eds.) ESORICS '10
29. Katz, J., Lindell, Y.: Introduction to Modern Cryptography. Chapman and Hall/CRC Press (2007)
30. Kopp, H., Bösch, C., Kargl, F.: Koppercoin - A distributed file storage with financial incentives. In: ISPEC'16
31. Kopp, H., Mödinger, D., Hauck, F.J., Kargl, F., Bösch, C.: Design of a privacy-preserving decentralized file storage with financial incentives. In: EuroS&P Workshops'17
32. Kosba, A.E., Miller, A., Shi, E., Wen, Z., Papamanthou, C.: Hawk: The blockchain model of cryptography and privacy-preserving smart contracts. In: S&P'16
33. Kupcu, A.: Efficient cryptography for the next generation secure cloud. Brown University (2010)
34. Kuppusamy, L., Rangasamy, J., Stebila, D., Boyd, C., Nieto, J.M.G.: Practical client puzzles in the standard model. In: Youm, H.Y., Won, Y. (eds.) ASIACCS '12
35. Labs, P.: Filecoin: A decentralized storage network (2017), https://filecoin.io/filecoin.pdf
36. Luu, L., Teutsch, J., Kulkarni, R., Saxena, P.: Demystifying incentives in the consensus computer. In: Ray, I., Li, N., Kruegel, C. (eds.) CCS'15
37. Ma, M.: Mitigating denial of service attacks with password puzzles. In: ITCC'05
38. Mahmoody, M., Moran, T., Vadhan, S.P.: Time-lock puzzles in the random oracle model. In: CRYPTO'11
39. Malavolta, G., Thyagarajan, S.A.K.: Homomorphic time-lock puzzles and applications. In: CRYPTO'19
40. Miller, A., Juels, A., Shi, E., Parno, B., Katz, J.: Permacoin: Repurposing bitcoin work for data preservation. In: S&P'14
41. Pedersen, T.P.: Non-interactive and information-theoretic secure verifiable secret sharing. In: CRYPTO '91
42. Renner, T., Müller, J., Kao, O.: Endolith: A blockchain-based framework to enhance data retention in cloud storages. In: PDP'18
43. Rivest, R.L., Shamir, A., Wagner, D.A.: Time-lock puzzles and timed-release crypto. Tech. rep. (1996)
44. Ruj, S., Rahman, M.S., Basu, A., Kiyomoto, S.: Blockstore: A secure decentralized storage framework on blockchain. In: AINA'18
45. Sengupta, B., Ruj, S.: Keyword-based delegable proofs of storage. In: ISPEC'18
46. Shacham, H., Waters, B.: Compact proofs of retrievability. In: ASIACRYPT. pp. 90–107 (2008)
47. Shen, S., Tzeng, W.: Delegable provable data possession for remote data in the clouds. In: ICICS 2011
48. Shi, E., Stefanov, E., Papamanthou, C.: Practical dynamic proofs of retrievability. In: CCS'13
49. Vorick, D., Champine, L.: Sia: Simple decentralized storage. Nebulous Inc (2014)
50. Waters, B., Juels, A., Halderman, J.A., Felten, E.W.: New client puzzle outsourcing techniques for dos resistance. In: CCS'04
51. Wesolowski, B.: Efficient verifiable delay functions. In: Ishai, Y., Rijmen, V. (eds.) EUROCRYPT'19
52. Wilkinson, S., Boshevski, T., Brandoff, J., Buterin, V.: Storj: a peer-to-peer cloud storage network (2014)
53. Xu, J., Yang, A., Zhou, J., Wong, D.S.: Lightweight delegatable proofs of storage. In: ESORICS'16

54. Xue, J., Xu, C., Bai, L.: Dstore: A distributed system for outsourced data storage and retrieval. Future Generation Comp. Syst. 2019
55. Xue, J., Xu, C., Zhang, Y., Bai, L.: Dstore: A distributed cloud storage system based on smart contracts and blockchain. In: ICA3PP'18
56. Yao, A.C.: Protocols for secure computations (extended abstract). In: FOCS'1982
57. Zhang, Y., Deng, R.H., Liu, X., Zheng, D.: Blockchain based efficient and robust fair payment for outsourcing services in cloud computing. Inf. Sci. (2018)

A Survey of Related Work

A.1 Time-lock Puzzles

The idea to send information into the *future*, i.e. time-lock puzzle/encryption, was first put forth by Timothy C. May. A time-lock puzzle allows a party to encrypt a message such that it cannot be decrypted until a certain amount of time has passed. In general, a time-lock scheme should allow generating (and verifying) a puzzle to take less time than solving it. The time-lock puzzle scheme that May proposed lies on a trusted agent. Later on, Rivest *et al* [43] propose a protocol that does not require a trusted agent and is secure against a receiver who may have access to many computation resources that can be run in parallel. It is based on Blum-Blum-Shub pseudorandom number generator that relies on modular repeated squaring, believed to be sequential. The scheme in [43] allows (only) the puzzle creator to verify the correctness of the puzzle solution using a secret key and the original secret message. This scheme has been the core of (almost) all later time-lock puzzles schemes that supports the encapsulation of an arbitrary message. Later on, [11, 21] proposed timed commitment schemes that offer more security properties, in the sense that they allow a puzzle generator to prove (in Zero-knowledge) to a puzzle solver that the correct solution (e.g. a signature of a public document) will be recovered after a certain time before the solver starts solving the puzzle. These schemes are more complex, due to the use of zero-knowledge proofs, and less efficient than [43]. Very recently, [39, 12] propose protocols for homomorphic time-lock puzzles, where an arbitrary function can be run over puzzles before they are solved. They mainly use fully homomorphic encryption and the RSA puzzle, proposed in [12], in a nutshell. In these protocols, all puzzles have an identical time parameter, and their solutions are supposed to be discovered at the same time. The main difference between the two protocols is the security assumption they rely on (i.e. the former uses a non-standard assumption while the latter relies on a standard one). Since both schemes use a generic fully homomorphic encryption, it is not hard to make them publicly verifiable. However, they are only of theoretical interest as in practice they impose high computation and communication costs, due to the use of fully homomorphic encryptions.

We also cover two related but different notions, pricing puzzles, and verifiable delay functions.

Pricing Puzzles. Also known as *client puzzles*. It was first put forth by Dwork *et al.* [17] who defined it as a function that requires a certain amount of computation resources to solve a puzzle. In general, the pricing puzzles are based on either hash inversion problems or number-theoretic. In the former category, a puzzle generator generates a puzzle as: $h = \text{H}(m||r)$, where H is a hash function, m is a public value and r is a

random value of a fixed size. Given $h, \mathtt{H}$ and m, the solver must find r such that the above equation holds. The size of r is picked in such a way that the expected time to find the solution is fixed (however it does not rule out finding the solution on the first attempt). The above hash-based scheme allows a solver to find a solution faster if it has more computational power resources running in parallel. The application area of such a puzzle includes defending against denial-of-service (DoS) attacks, reaching a consensus in cryptocurrencies, etc. A variant of such a puzzle uses iterative hashing; for instance, to generate a set of puzzles in the case where the solver receives a service proportional to the number of puzzles it solves [24], or to generate password puzzle to mitigate DoS attacks [37]. However, the iterative hashing schemes are partially parallelizable, in the sense that every single invocation of the hash function can be run in parallel. Later on, [38] investigates the possibility of constructing (time-lock) puzzles in the random oracle model. Their main result was negative, that rules out time-lock puzzles that require more parallel time to solve than the total work required to generate. Also [38] proposes an iterative hash-based mechanism (very similar to [37]) that allows a puzzle generator to generate a puzzle with n parallel queries to the random oracle, but the solver needs n rounds of serial queries. Nevertheless, this scheme is also partially parallelizable, as each instance of the puzzle can be solved in parallel. Note that the above hash-based puzzle schemes would have very limited applications if they are used directly to encapsulate a message: m' of arbitrary size. The reason is that, in these schemes, the solution size: $|r|$ plays a vital role in (adjusting) the time taken to solve the puzzle. If the solution size becomes bigger, as a result of combining r with m', i.e. $r \odot m'$, then it would take longer to find the solution. This means the puzzles can be used only in the cases where the time required to find a solution is long enough, and is a function of $|r \odot m'|$, which seriously restricts its application. Researchers also propose non-parallelizable pricing puzzles based on number theoretic [50, 34, 28] whose main application is to resist DoS attacks. These schemes have a more efficient verification mechanism than the one proposed in [43]. But, they are only privately verifiable and not designed to encapsulate an arbitrary message.

Verifiable Delay Function (VDF). Allows a prover to provide a publicly verifiable proof stating it has performed a pre-determined number of sequential computations. It has many applications, e.g. in decentralised systems to extract trustworthy public randomness from a blockchain. VDF was first formalised by Boneh *et al* in [10] who proposed several VDF constructions based on SNARKs along with either incrementally verifiable computation or injective polynomials, or based on time-lock puzzles, where the SNARKs based approaches require a trusted setup. Later on, [51] improved the previous VDF's from different perspectives and proposed a scheme based on RSA time-lock encryption, in the random oracle model. To date, this protocol is the most efficient VDF. It also supports batch verification, such that given a single proof a verifier can efficiently check the validity of multiple outputs of the verifiable delay function. As discussed above, (most of) VDF schemes are built upon time-lock puzzles, however the converse is not necessarily the case, as VDF's are not designed to encapsulate an arbitrary private message, and they take a public message as input while time-lock puzzles are designed to conceal a private input message.

A.2 Proof of Storage

Traditional Proof of Storage Proof of storage (PoS) is a cryptographic protocol that allows a client to efficiently verify the integrity or availability of its data stored in a remote server, not necessarily trusted [27]. PoS can be categorised into two broad classes: Proofs of Retrievability (PoR) and Proofs of Data Possession (PDP). The main difference between the two categories is the level of security assurance provided. PoR schemes guarantee that the server maintains knowledge of *all* of the client's outsourced data, while PDP protocols only ensure that the server is storing *most* of the client data. From a technical point of view, the main difference in most prior PDP and PoR constructions is that PoR schemes store a redundant encoding of the client data on the server by employing an error-correcting code, e.g. Reed-Solomon codes, while such encoding is not utilised in PDP schemes. Hence, PoR schemes provide stronger security guarantees compared to PDP at the cost of additional storage space and encoding/decoding, (for more details see [33, 48, 15]). Furthermore, each PoR and PDP scheme can be also grouped into two categories; namely, publicly and privately verifiable. In a publicly verifiable scheme, everyone without knowing a secret can verify proof, whereas a verifier in a privately verifiable scheme requires the knowledge of a secret.

The notion of PoR first was introduced and defined in [26], where the authors designed a protocol that utilises random sentinels, symmetric key encryption, error-correcting code, and pseudorandom permutation. In this protocol, a client in the setup phase, applies error-correcting code to every file block and encrypts each encoded block. Then, it computes a set of random sentinels and appends them to the encrypted file. It permutes all values (i.e. sentinels and encrypted file blocks) and sends them to the cloud server. To check if the server has retained the file, the client specifies random positions of some sentinels in the encoded file and asks the server to return those sentinel values. Next, the client checks if it gets the sentinels it asked for. In this scheme, the security holds, as the server cannot feasibly distinguish between sentinels and the actual file blocks and the sentinels have been distributed uniformly among the file blocks. But, as individual sentinel is only one-time verifiable, there is an upper bound on the number of verifications performed by the client, and when reached, the client has to re-encode the file. To overcome the issues related to the bounded number of verifications, the authors have also suggested that sentinels can be replaced with MAC on every file block or a Merkle tree constructed on the file blocks. This protocol is computationally efficient and its communication cost is linear with the number of challenges sent[4]. The sentinel and MAC-based schemes above are privately verifiable while the one uses a Merkle tree supports public verifiability. However, the publicly verifiable Merkle tree-based approach has a communication cost logarithmic with the file size and the prover has to send a set of file blocks to the verifier that yields a high communication cost.

Later on, [46] improves the previous sentinel-based scheme and definition, and proposes two PoR protocols (with an unlimited number of verifications), one utilises MAC and the other one relies on BLS signatures. In particular, a client at setup phase, encodes

---

[4] It is not hard to make the communication cost of this scheme constant; for instance, by letting the server order the challenged sentinels, concatenate them and send a hash of the concatenated value to the client.

its file blocks using error-correcting code and then for each encoded file block it generates a tag that can be either a MAC or BLS signature of that block. In the verification phase, it specifies a set of random indices corresponding to the file's blocks; and accordingly the server sends a proof to it. These two schemes have a low communication cost, as due to the homomorphic property of the tags, the prover can aggregate proofs into a single authenticator value and no file blocks are sent to the prover. Also, the protocol based on MAC supports efficient private verification. But, the one that uses BLS signatures supports public verifiability at the cost of public key operations and it is far less efficient than the MAC-based one. Within the last decade, researchers have extended PoR protocols from several perspectives, e.g. those that support: efficient update [48], the delegation of file pre-processing [2], or the delegation of verification [3].

On the other hand, PDP was first introduced in [4]. PDP protocols focus only on verifying the integrity of outsourced data. In this setting, clients can ensure that a certain percentage of data blocks are available. The authors in [4] propose two schemes, for public and private verification both of which utilise RSA-based homomorphic verifiable tags generated for each file block and use the spot-checking techniques (similar to [46]) to check a random subset of a file's blocks. In both schemes, the proof size is constant, while the verification cost is high as it requires many exponentiations over an RSA ring. Later on, an efficient and scalable PDP scheme that supports a limited number of verifications is proposed in [6]. The scheme supports only privately verifiable PDP and is based on a combination of pre-computation technique and symmetric key primitives. Ever since, researcher proposed different variants of PDP, e.g. dynamic PDP [18], multi-replica PDP[19], keyword-based PDP [45].

Thus, (a) in publicly verifiable PoS it is assumed the verifier is fully trusted with the verification correctness, (b) these schemes either have a very high verification cost when the proof size is constant (when signature-based tags are used) or have a communication cost logarithmic with the entire file size if their verification is efficient (when a Merkle tree is utilised), and (c) the privately verifiable proofs can have a constant proof size and efficient verification algorithm, but the data owner has to perform the verification.

Outsourced PoR Recently, [3, 53] propose *outsourced* PoR protocols that allow clients to outsource the verification to a third party auditor not necessarily trusted. The scheme in [3] uses MAC-based tags, zero-knowledge proofs, and error-correcting codes. At a high level, the protocol works as follows. At the setup, the client encodes its data using error-correcting codes, generates MAC on every file block, and stores the encoded file and tags on the server. Then, the auditor downloads the encoded file, generates another set of MAC's on the file blocks. It uploads the MAC's to the cloud. Also, the auditor proves to the client in zero-knowledge that it has created each MAC correctly. To do that, it uses a non-interactive zero-knowledge proof in the random oracle model. If the client accepts all zero-knowledge proofs, it sings every proof and sends the signatures back to the auditor. In the verification phase, the auditor sends two sets of random challenges extracted from a blockchain. Upon receiving the challenges, the server provides two separate proofs, one for the auditor and the other one for the client. The auditor verifies the proof generated for it and locally stores the clients' proofs. In the case where the auditor's proof is not accepted, an honest auditor would inform the client who will

come online and checks both its proofs and the auditors' proof. The scheme provides two layers of verification to the client: *CheckLog* and *ProveLog*. *CheckLog* is more efficient than the other one, as the auditor sends much fewer challenges to the server to generate the client's proof for each verification. In the case where the client's check fails, the client will assume that either the server or auditor has acted maliciously, and to pinpoint the malicious party, it needs to proceed to *ProveLog* where allows the client to audit the auditor. In this verification, the auditor must reveal its secret keys with which the client checks all the proofs provided by the server to the auditor for the entire period that the client was offline. In this protocol, the verification phase is efficient (due to the use of MACs). In particular, the verification's computation cost for the auditor and client is linear with the number of challenges and their communication cost is constant in the file size.

Although the protocol in [3] is appealing, it has several shortcomings: (a) auditor can get a free ride: in the case where a known highly reputable server, e.g. Google or Amazon, always generates accepting proofs for both client and auditor, an economically rational auditor can skip performing its part (e.g. to save computation cost) and get still paid by the client. This deviation from the protocol can never be detected by the client in the protocol. (b) no guarantee for a real-time detection/notification: the client may not be notified in the real-time about the data unavailability if the auditor is malicious; therefore, this scheme is not suitable for the case where a client's involvement for the data extraction is immediately needed once a misbehaviour is detected, e.g., in the case of hardware depreciation. (c) a high cost of auditor onboarding: revoking an auditor and onboarding a new one, imposes a high cost on the client and new auditor, as new auditor need to re-run the setup phase (that includes data downloading, zero-knowledge proofs) and the client needs to verify and sign the zero-knowledge proofs. (d) *ProveLog* costs even an honest auditor: the only way for the client to ensure the auditor has fully followed the protocol is to run *ProveLog* that requires the auditor to reveal its secret values. After that, an honest auditor has to re-run the computationally expensive setup again. Since the auditor is not trusted and no guarantee that it alerts the client as soon as the server's misbehaviour is detected, its involvement in the protocol seems unnecessary; and the whole scheme can be replaced with a much simpler one: the client, similar to a standard privately verifiable PoR, encodes its data and stores it in the server. Then, for each verification time, the server gets the random challenges from the blockchain and publishes proofs in bulletin board (to timestamp it). When the client comes online, it checks the proofs and accordingly takes the required action, e.g. pays the server, tries to recover the file.

Xu et al. in [53] propose a publicly verifiable PoR to improve the computation cost at setup. The scheme uses BLS signatures-like tags, polynomial arithmetics, and polynomial commitments; unlike previous schemes, each tag has a more complex structure. In this protocol, an auditor is assumed to be fully trusted during the verifications. Later on, however, when it is revoked by the client it may become malicious, i.e. may collude with the server and reveal its secret. This scheme allows a client to update its tags when a new auditor joins. To do that, the client needs to download all the blocks' tags, refresh them locally, and uploaded the fresh tags to the server. The verification for auditor and client involves public key operations and is more expensive than the one in [3]. The use

of an auditor's revocation remains unclear in the paper, as auditors are assumed to be honest in this protocol. A delegatable PDP is proposed in [47] to ensure only authorised parties can verify the integrity of remote data. However, the delegated party in this scheme is fully trusted with the correctness of verification it performs. The scheme uses authenticator tags similar to BLS signatures based ones.

Very recently, proof of "storage-time" has been proposed in [5]. The paper proposes two protocols: basic PoSt and compact PoSt. At a high level, they offer the guarantee to a client that its data has been available on a remote storage server for a fixed time period: T. The idea is that a client uploads to a server its data once, then the server generates proofs of storage (e.g. PoR) periodically, collects them and sends the collection after the time T to a verifier (i.e. client in basic PoSt or third-party in compact PoSt) who can check the proof. The first protocol, basic PoSt, uses Merkle tree-based PoR and VDF. In this protocol, The client precomputes a set of metadata (tags, challenges, and their related proofs). It sends the file, metadata, and a single challenge to the server. The server uses the challenge to generate a PoR. It feeds the hash of the PoR to VDF to generate an output. It considers the hash of VDF's output as the next challenge from which the next PoR is generated. This process goes on until all z PoR's are generated. It sends all PoR proofs along with the proofs proving the correctness of VDF's outputs to a verifier. Given the two sets of proofs, a verifier can check their correctness and ultimately conclude that the file was retrievable within the time period. The scheme is publicly verifiable. However, the verification's computation cost is significant. As, it requires the validator to validate the correctness of VDF's outputs, that imposes a high cost. In total (for z PoR proofs) it involves at least $3z$ modular exponentiations even if it uses the fastest VDF, proposed in [51]. Such costs are missed out and not taken into account in the paper. Moreover, the use of the Merkle tree introduces a high communication cost, as well. The authors suggest a smart contract can play the validator's role. However, we argue that this would impose a significant financial cost to the contract and users due to very high computation and communication costs stem from the verification and prove algorithms. The second protocol, compact PoSt, allows the server to combine PoR proofs that ultimately reduces the communication cost. The protocol's design is very similar to the previous one, but it replaces the Merkle tree approach with simply hashing the entire file and replaces the VDF with a trapdoor delay function (TDF) that requires a secret to generate/verify the delay function's output. Therefore, this protocol is only privately verifiable. The paper claims that the verification in this protocol can be performed by a trusted third-party, e.g. smart contract. Nevertheless, we argue that this is not the case. As, the scheme requires a set of secrets (e.g. challenges) to perform the verification, while a smart contract does not maintain a private state. If the challenges are given to the contract, then the server can read the contract and compute all proofs in one go (which violates the security requirement set out in the paper). Moreover, the same security issue would arise in the case where the verification is delegated to a third-party auditor who may collude with the server (as the auditor can send all challenges to the server at once). Therefore, the only secure option, in the compact PoSt, is that the verifier performs the validation itself (i.e. private verification/validation). Note that all the above outsourced PoR schemes [3, 53, 5] assume the client behaves honestly to-

wards the server. Otherwise, a malicious client can generate the tags in a way that makes an honest server generate invalid proofs.

Hence, the existing outsourced PoR protocols either suffer from several security issues and guarantee no real-time detection or have to fully trust verifiers with the verification correctness, or are very inefficient.

Distributed PoR using a (Tailored) Blockchain Distributed PoR allows data to be distributed to numerous storage servers to achieve robustness, and address a single point of failure issue. Permacoin [40] is one of those that distribute data among customised blockchain nodes and repurposes mining resources of Bitcoin blockchain miners. In Permacoin, each miner needs to prove that it has a portion of the file and to do so it provides proof of data retrievability verified by other miners. The miner, in this scheme needs to invest on both computation resources and storage resources (to generate proof of retrievability). The protocol uses a Merkle tree built on top of file blocks to support publicly verifiable PoR. The mining procedure in this protocol is based on iterative hashing. In Permacoin, in each verification, the prover has to send both challenged file blocks and the proof path in the tree, that imposes communication cost logarithmic to the size of the *entire original file*. In this scheme, an accepting proof only indicates a portion of the data holding by that miner (prover) is retrievable. Thus, for a data owner to have a guarantee that the entire data is retrievable, it has to either wait for a long period of time (depending on the file size and the number of miners) or hope that there are enough active miners in each epoch. Also, since the miners are both verifiers and provers (storage providers), it is assumed that the storage providers are economically rational (which is weaker than a malicious one).

Filecoin and KopperCoin in [35, 30] respectively, offer similar features, i.e. repurposing Bitcoin and supporting distributed PoR. However, Filecoin, in addition to a Merkle tree, utilises generic zero-knowledge proofs (i.e. zk-SNARKS) that result in a high overall computation cost and requires a trusted party to generate the system parameters. Filecoin uses proof of work as well as PoR. On the other hand, KopperCoin uses BLS signature-based publicly verifiable PoR that has a constant communication cost but has a high computation cost. Unlike Filecoin, KopperCoin does not use a PoW.

In the same line of research, [31] proposes a customised blockchain that supports distributed PoR as well as preserving the privacy of on-chain payments between storage users and providers. It however is computationally expensive as it uses publicly verifiable tags based on BLS signatures for PoR and ring signatures for the privacy-preserving payments. Likewise, [44] proposes a high-level distributed PoR framework whose aim is efficient utilization of storage resources offered by storage providers. It also uses BLS signatures and a Merkle tree along with a smart contract for payments. Similarly, [55, 54] propose schemes that allow a user to store its encrypted data in the blockchain. The scheme uses a Merkle tree for PoR and a smart contract to transfer fees to the blockchain nodes storing the data. In these two schemes, the data owner is the party who sends random challenges to storage nodes, so it has to be online when verification is needed. Storj [52] also falls in this category where there is a tailored blockchain comprising a set of storage nodes that store a part of data and provide proofs when they

are challenged. In Storj, there are trusted nodes, Satellite, who do the verifications on behalf of the clients. The scheme uses a Merkle tree-based PoR.

So, existing distributed PoR protocols have either a large proof size or high verification cost. Also, they do not guarantee that the *entire* file is retrievable in *real-time*.

Verifying Remote (off-chain) data via a Blockchain To relax the assumption that an auditor is fully trusted with the correctness of verification in the publicly verifiable PoR schemes while storing data off-chain, e.g. in a cloud, researchers proposed numerous protocols, e.g. [42, 25, 57, 20, 8, 49], that delegate the verification procedure to blockchain nodes. The high-level framework proposed in [42] requires only a hash of the entire file is stored in a smart contract, where when later on the user access the file, it computes the hash of the file and compares it with the one stored in the contract. In this scheme, the entire file has to be accessed by the user for each verification which is what exactly PoR schemes avoid doing. The protocol in [25] uses BLS signature-based tags and Merkle tree where the tags are broadcast to all blockchain nodes. However, surprisingly, the protocol assumes the cloud server is fully trusted and stores data safely, which raises the question that *"why is a data verification needed in the first place?"*. In this protocol, the tags are generated by the cloud who sends them to the nodes. The tags are never checked against the outsourced data. So, the only security guarantee [25] offers is the immutability of tags. The high-level scheme proposed in [57] allows a client to pay the storage server in a fair manner. The scheme uses a blockchain for payment and Merkle tree for PoR. In this protocol, the client has to be online and send random challenges to the server for every verification.

Audita [20] utilises an augmented blockchain and RSA signature-based PDP [4] to achieve its goal. In this scheme miners, for each epoch, pick a dealer who sends a challenge to the storage node(s) to get proofs of data possession. Similar to Permacoin [40], Audita substitutes proof of work with PDP, so if a proof is accepted a new block is added to the chain. But, in Audita every miner carries out expensive public key based verifications. The protocol proposed in [8], similar to [3], uses a third-party auditor. It mainly utilises, a smart contract and BLS signature-based tags. In this protocol, unlike [3], when the auditor raises a dispute during the verification it calls a smart contract who performs the verification again to detect a misbehaving party, i.e. the server or auditor. The protocol is computationally more expensive than [3] and inherits the same issues, i.e. issues (a-c) stated above. Sia [49] is a mechanism in which a data owner distributes its data among off-chain storage servers who periodically provide a PoR to a smart contract signed between the data owner and the servers. Each server gets paid if its proof is accepted by the contract. In this scheme, a Merkle tree-based PoR is used.

As evident, the schemes designed for blockchain-based verification of data stored in a storage server either require clients to access whole outsourced data for every verification, or impose a high communication/computation cost, or clients have to be online for each verification.

Fair Exchange of Digital Services Campanelli *et al.* [13] propose a scheme that allows different parties to exchange digital services (and goods) over Bitcoin blockchain. This scheme, for instance in PoR context, allows the storage provider to get paid if and

only if the data owner receives an accepting proof. It has two variants, publicly and privately verifiable both of which use a smart contract. In addition, in the privately verifiable variant, MAC-based tags and generic secure multi-party computation are used, e.g. Yao's garbled circuit [56], while in the publicly verifiable one BLS signature tags and zk-SNARK are utilised. Nevertheless, this scheme assumes that either the client is available and online when PoR is provided (in privately verifiable variant) or the third party, acting on a client's behalf, performs PoR verification honestly (in publicly verifiable one).

B   Notations

We summarise our notations in Table 4.

Table 4: Notation Table.

Setting	Symbol	Description	Setting	Symbol	Description
Generic	z	Number of puzzles or delegated verifications	**SO-PoR**	PRF	Pseudorandom function
	h, h_j	Hash values		$\hat{k}, v_j, l_j$	PRF's keys
	d, d_j	Randomness of commitment		ι	Security parameter, $\iota = 128$-bit
	n	Number of file blocks		p	Large prime number, $|p| = \iota$
	m, m_j	Plaintext messages		w	Blockchain block index
	$\ddot{o}$	Pair representation		g	Blockchain security parameter: chain quality
	$\vec{\sigma}$	Vector representation		λ'	Blockchain generic security parameter
	$\ddot{p} : (m_j, d_j)$	Commitment opening		$\boldsymbol{F}$	Outsourced encoded file
	$\mathtt{H}$	Hash function		F_j	A file block
C-TLP	λ	TLP security parameter		$|\boldsymbol{F}|$	Number of file blocks, $|\boldsymbol{F}| = n$
	Δ	Time win. message remains hidden		$||\boldsymbol{F}||$	File bit-size
	S	Max. squaring done per sec.		σ_i	Permanent tag
	T	$T = S\Delta$		$\sigma_{b,j}$	Disposable tag
	s_j	j-th solution		$\alpha, \alpha_j, r_i, r_{b,j}$	Pseudorandom values
	$\ddot{o}_j$	j-th puzzle, $\ddot{o}_j : (o_{j,1}, o_{j,2})$		c	Number of challenges
	f_j	Time when j-th solution is found		$\mathcal{B}_j$	Blockchain's j-th block
	k, k_j	Sym. key encryption keys		(μ_j, ξ_j)	j-th PoR proof
	pk, sk	Public and secret keys		Δ_1	Time taken to generate a PoR
	q_1, q_2	Large prime numbers		Δ_2	Time taken a contract gets a message
	N	RSA modulus, $N = q_1 q_2$		e	Coins paid for an accepting PoR

C   Discussion on Time-lock Puzzle

In general, time-lock puzzles by definition are sequential functions. In their construction, it is required a function inherently sequential, e.g. modular squaring, is called iteratively for a certain number of times. The notions of sequential function and iterated sequential functions, in the presence of an adversary possessing a polynomial number of processors, have been generalised and defined in [10]. Below, for the sake of completeness, we provide those definitions.

Definition 8 ($\Delta, \delta(\Delta)$)**-Sequential function**)**.** *For a function:* $\delta(\Delta)$, *time parameter:* Δ *and security parameter:* $\lambda = O(\log(|X|))$, $f : X \rightarrow Y$ *is a* ($\Delta, \delta(\Delta)$)*-sequential function if the following conditions hold:*

- *There exists an algorithm that for all* $x \in X$ *evaluates* f *in parallel time* Δ *using* $poly(\log(\Delta), \lambda)$ *processors.*

- *For all adversaries $\mathcal{A}$ that run in parallel time strictly less than $\delta(\Delta)$ with $poly(\Delta, \lambda)$ processors:*

$$Pr\left[y_A = f(x) \middle| y_A \xleftarrow{\$} \mathcal{A}(\lambda, x), x \xleftarrow{\$} X\right] \leq negl(\lambda)$$

where $\delta(\Delta) = (1 - \epsilon)\Delta$ and $\epsilon < 1$

Definition 9 (Iterated Sequential function). *Let $g : X \to X$ be a $(\Delta, \delta(\Delta))$-sequential function. A function $f : \mathbb{N} \times X \to X$ defined as $f(k, x) = g^{(k)}(x) = \underbrace{g \circ g \circ \ldots \circ g}_{k \times}$ is an iterated sequential function, with round function g, if for all $k = 2^{o(\lambda)}$ the function $h : X \to X$ defined by $h(x) = f(k, x)$ is $(k\Delta, \delta(\Delta))$-sequential.*

The main property of an iterated sequential function is that iteration of the round function g, is the fastest way to evaluate the function. Iterated squaring in a finite group of unknown order, is widely believed to be a candidate for an iterated sequential function. Its definition is as follows.

Assumption 1 (Iterated Squaring) *Let N be a strong RSA modulus, r be a generator of $\mathbb{Z}_N$, Δ be a time parameter, and $T = poly(\Delta, \lambda)$. For any $\mathcal{A}$, defined above, there is a negligible function $\mu(.)$ such that:*

$$Pr\left[\mathcal{A}(N, r, y) \to b \middle| \begin{array}{l} r \xleftarrow{\$} \mathbb{Z}_N, b \xleftarrow{\$} \{0, 1\} \\ if\ b = 0,\ y \xleftarrow{\$} \mathbb{Z}_N \\ otherwise\ y = r^{2^T} \end{array}\right] \leq \frac{1}{2} + \mu(\lambda)$$

Theorem 3 (RSA-based TLP Security). *Let N be a strong RSA modulus and Δ be the period within which the solution/secret remains hidden. If the sequential squaring assumption holds, factoring N is a hard problem and the symmetric key encryption is semantically secure, then RSA-based TLP scheme (presented in Section 3.4) is a secure time-lock puzzle.*

Proof (Proof sketch). Let a solver be $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$, where $\mathcal{A}_1$ runs in total time $O(poly(\Delta, \lambda))$ and $\mathcal{A}_2$ runs in time $\delta(\Delta) < \Delta$ using at most $\pi(\Delta)$ parallel processors. For the solver to find the secret significantly earlier than $\delta(\Delta)$, given the public parameters, it needs to either: (a) compute $\phi(N)$, so it can generate the blinding factor: b as fast as it is done in the encryption phase, or (b) break the symmetric key encryption, or (c) extract k from o_2 by finding the blinding factor without performing a sufficient number of squaring (and without knowing $\phi(N)$). However, finding $\phi(N)$ is as hard as factoring N, also as long as the encryption is secure and k is sufficiently large it cannot break the symmetric key encryption. Moreover, the blinding factor is a uniformly random element of the ring (due to Assumption 1), so it prevents the solver from finding the blinding factor without carrying out enough squaring within time significantly less than $\delta(\Delta)$. Thus, any adversary $\mathcal{A}$ cannot find the secret without carrying out a sufficient number of squaring that would take it $\delta(\Delta)$ □

D   Further Remarks on C-TLP Protocol

In this section, we provide some remarks on C-TLP scheme presented in Section 5.3.

Remark 3. Recall, to make each puzzle instance, a distinct random generator: r_j, is used. This is the reason, in the C-TLP protocol, before a puzzle is generated in step 2b, a new public key is set in step 2a. Also, at the beginning of the protocol only r_1 is public and the rest of the generators are kept secret. They are found and used sequentially after their related puzzle is solved.

Remark 4. The commitments opening, including the commitment random values, are not known to other verifiers (than the puzzle generator) at the beginning of the protocol. At this point, only the committed values are public. Once a solver solves each puzzle, it extracts one of the commitments' opening and sends it to a public verifier who can check if the opening matches the commitment.

Remark 5. In C-TLP, we use the folklore hash-based commitment scheme, in the random oracle model, only to achieve more computation improvement than that can be achieved in the standard model. But C-TLP can utilise any efficient non-interactive commitment scheme in the *standard model* as well, e.g. Pedersen Commitment.

Remark 6. The efficiency of C-TLP scheme stems from three crucial factors: (a) removing computation overlaps when solving different puzzles: even though solving j-th puzzle, where $j > 1$, requires jT squaring, $(j-1)T$ of the squaring is used to solve previous puzzles that leads to $\frac{z+1}{2}$ times computation cost reduction at the server-side, (b) supporting reusable single public parameter: $a = 2^T$, generated only once that costs $O(1)$, as opposed to the RSA-based TLP whose cost is linear: $O(z)$, and (c) supporting efficient verification: due to the way each message is encoded (i.e. embedding the opening in a solution).

Remark 7. C-TLP also can efficiently be used in a multi-server setting, where there are z servers: $\{S_1, ..., S_z\}$, each S_j needs to solve puzzle $\ddot{o}_j$ at time f_j and passes on the solution to the next server S_{j+1} to solve the next puzzle by time $f_{j+1} > f_j$. In this setting, due to the scalability property of C-TLP (and unlike using the existing time-lock puzzles naively), other servers do not need to start solving the puzzle as soon as the client releases puzzles' public parameters. Instead, they can wait until the previous solution is issued that saves them a significant cost. Furthermore, a server can first verify the correctness of the solution found by the previous server (due to the public verifiability of C-TLP), if accepted then it starts finding the next solution.

Remark 8. In the following, we outline an approach that looks an option to construct an efficient C-TLP; however, as we will show it would not be secure. In particular, one uses the TLP to generate z public and secret key pairs. Then, it uses the TLP to compute z-th puzzle as $\texttt{TLP.GenPuZ}(m_z, pk_z, sk_z) \rightarrow \ddot{o}_z$. Then, it embeds $\ddot{o}_z$ into $(z-1)$-th one, i.e. $\texttt{TLP.GenPuZ}(m_{z-1}||\ddot{o}_z, pk_{z-1}, sk_{z-1}) \rightarrow \ddot{o}_{z-1}$. This process goes on until $\ddot{o}_1$ is created. It sends the combined puzzles and public key (including all random generators) to the server; with the hope that puzzles can be solved sequentially and the time gap between finding two solutions will be Δ. This approach is not secure, because as soon as the server accesses $\ddot{o}_1$ and public parameters, it can in parallel perform T squaring on every generator, i.e. $r_i^{2^T}$, for all $i, 1 \leq i \leq z$. In this case, as soon as $\ddot{o}_1$ is solved and $\ddot{o}_2$ is extracted, it has enough information to immediately solve $\ddot{o}_2$ and accordingly the rest of the puzzles without doing any further exponentiation.

E C-TLP Security Proof

In this section, we present the security proof of C-TLP scheme. We first prove that without solving j-th puzzle, a solver cannot find the parameters needed to solve the next puzzle, i.e. $(j + 1)$-th one.

Lemma 1 (Next Group Generator Privacy). *Let k be a random key for a symmetric key encryption, and N be a sufficiently large RSA modulus. Let the security parameter be $\lambda = |N| = |k|$. In C-TLP, given puzzle vector: $\vec{o}$ and public key: pk, an adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$, defined in Section 4.3, cannot find the next group generator: r_{j+1}, where $r_{j+1} \xleftarrow{\$} \mathbb{Z}_N^*$ and $j \geq 1$, significantly smaller than $T_j = \delta(j\Delta)$, except with a negligible probability in the security parameter, $\mu(\lambda)$*

Proof. Since the next generator: r_{j+1}, is: (a) encrypted along with the j-th puzzle solution: s_j, and (b) picked uniformly at random from $\mathbb{Z}_N^*$, for the adversary to find r_{j+1} without performing enough squaring, i.e. T_j, it has to either (a) break the symmetric key scheme, decrypt the related ciphertext: s_i and extract the random value from it, or (b) correctly guess r_{j+1}. In both cases, the probability of success is negligible in secure parameter $\mu(\lambda)$, i.e. $2^{-|k|}$ in the former case and $2^{-|N|}$ in the latter one. □

In the following, we prove that the privacy of a solution in C-TLP scheme is preserved according to Definition 5.

Theorem 4 (C-TLP Solution Privacy). *Let N be a strong RSA modulus and Δ be a time parameter. If the sequential squaring assumption holds, factoring N is a hard problem, $\mathtt{H}(.)$ is a random oracle and the symmetric key encryption is semantically secure, then C-TLP encoding z solutions is a privacy-preserving multi-instance time-lock puzzle w.r.t. Definition 5.*

Proof. In the following, we argue for an adversary $\mathcal{A} = (\mathcal{A}_1, \mathcal{A}_2)$, where $\mathcal{A}_1$ runs in total time $O(poly(j\Delta, \lambda))$, $\mathcal{A}_2$ runs in time $\delta(j\Delta) < j\Delta$ using at most $\pi(\Delta)$ parallel processors, and $j \in [1, z]$, (a) when $z = 1$: to find s_1 earlier than $\delta(\Delta)$, it has to break the TLP scheme, and (b) when $z > 1$: to find s_j earlier than $T_j = \delta(j\Delta)$, it has to either find at least one of the previous solutions earlier than it is supposed to (that ultimately requires breaking TLP scheme again), or find j-th generator: r_j, earlier. Also, we argue that the commitments: h_j, are computationally hiding. Specifically, when $z = 1$, the security of C-TLP is reduced to the security of the TLP and the scheme is secure as long as TLP is, as the two schemes would be identical. On the other hand, when $z > 1$, the adversary has to either find s_j earlier than T_j as soon as the previous solution: s_{j-1} is found that requires either breaking the TLP scheme, or finding any generator r_j before s_{j-1} is extracted, when $j \in [2, z]$. Nevertheless, the TLP scheme is secure (under RSA, sequential squaring, and security of symmetric key encryption assumptions) according to Theorem 3, and also the probability of finding the next generator: r_j earlier than T_{j-1} is negligible, according to Lemma 1. Moreover, for an adversary to find a solution earlier, it may also try to find a (partial information of) pre-image of the commitment: h_j before fully (or without) solving the puzzle. But, this is infeasible for a PPT adversary, given output of a random oracle: $\mathtt{H}(.)$. Thus, C-TLP is a privacy-preserving multi-instance time-lock puzzle scheme. □

Next, we prove that the validity of a solution in C-TLP scheme is preserved according to Definition 6.

Theorem 5 (C-TLP Solution Validity). *Let* $\mathtt{H}(.)$ *be a hash function modeled as a random oracle. Then, C-TLP preserves a solution validity w.r.t. Definition 6.*

Proof. The proof boils down to the security (i.e. binding property) of the traditional hash-based commitment scheme. In particular, given an opening pair, $\ddot{p} : (m_j, d_j)$ and the commitment $h_j = \mathtt{H}(m_j, d_j)$, for an adversary to break the solution validity, it has to come up (m'_j, d'_j), such that $\mathtt{H}(m'_j, d'_j) = h_j$, where $m_j \neq m'_j$, i.e. finds a collision of $\mathtt{H}(.)$. However, this is infeasible for a PPT adversary, as $\mathtt{H}(.)$ is collision resistance, in the random oracle model. □

In the following, we restate the main theorem presented in Section 4.4 and then prove it.

Theorem 1 (C-TLP Security). *C-TLP is a secure multi-instance time-lock puzzle.*

Proof. According to Theorems 4 and 5, the privacy and validity of a solution in C-TLP are preserved, respectively. So, w.r.t. Definition 7, C-TLP is a secure multi-instance time-lock puzzle. □

F   Asymptotic Cost Analysis of C-TLP

In this section, we analyse the communication and computation complexity of C-TLP. We consider a generic setting where the protocol deals with z puzzles.

Computation Complexity. For a client to generate z puzzles, in total: in step 1a, it performs one exponentiation over $\mathrm{mod}\,\phi(N)$. In step 2b, it z times calls $\mathtt{TLP.GenPuZ}(.)$. This in total involves z symmetric key-based encryption, z modular exponentiations over $\mathbb{Z}_N$ and z modular additions. Also, in step 2c it performs z invocations of a commitment scheme (to commit), i.e. if a hash-based commitment is used then it would involve z invocations of a hash function, and if Pedersen commitment is used then it would involve $2z$ exponentiations and z multiplications, where all operations, in the latter commitment, are done over a $\mathrm{mod}\,q$ for a large prime number: q, e.g. $|q| = 1024$-bit. Thus, the overall computation complexity of the client is $O(z)$. For the server to solve z puzzles, it z times calls $\mathtt{TLP.SolvPuz}(.)$. This in total involves Tz modular squaring over $\mathbb{Z}_N$, z modular additions over $\mathbb{Z}_N$, and z symmetric key based decryption. The server's cost of proving, in step 4, is very low, as it involves only parsing z strings. Therefore, the server total computation complexity is $O(Tz)$. The verification cost, in step 5, only involves z invocation of commitment scheme (to verify each opening) and it is *independent* of the RSA security parameters. If the hash-based commitment is used, then it would involve z invocations of a hash function, if Pedersen commitment is utilised, then in total it would involve $2z$ exponentiations and z multiplications performed over $\mathrm{mod}\,q$. Thus, the verification's complexity is $O(z)$.

Communication Complexity. In step 2, the client publishes two vectors: $\vec{o}$ and $\vec{h}$, with $2z$ and z elements respectively. Each element of $\vec{o}$ is a pair $(o_{j,1}, o_{j,2})$, where $o_{j,1}$ is an

output of symmetric key encryption, e.g. $|o_{j,1}| = 128$-bit, and $o_{j,2}$ is an element of $\mathbb{Z}_N$, e.g. $|o_{j,2}| = 2048$-bit. Also, each element h_j of $\vec{h}$ is either an output of a hash function, when a hash-based commitment is used, e.g. $|h_j| = 256$-bit, or an element of $\mathbb{F}_q$ when Pedersen commitment is used, e.g. $|h_j| = 1024$-bit. Therefore, its total bandwidth is about $2432z$ bits when the former, or $3200z$ bits when the latter commitment scheme is utilised. Also, its complexity is $O(z)$. Note, in C-TLP, when a server finds a solution, it does not broadcast anything, instead it moves on to the next step: $\texttt{Prove}(.)$. For the server to prove, in step 4, in total, it sends z pairs (m_j, d_j) to the verifier, where m_j is an arbitrary message, e.g. $|m_j| = 500$-bit, and d_j is either a long enough random value, e.g. $|d_j| = 128$-bit, when the hash-based commitment is used, or an element of $\mathbb{F}_q$ when Pedersen scheme is used, e.g. $|d_j| = 1024$-bit. Therefore, its bandwidth is about either $628z$ or $1524z$ bits when the former or latter commitment scheme is used respectively. The solver's communication complexity is $O(z)$.

G Traditional PoR Model

In this section, we restate a definition of the traditional PoR scheme [46]. In general, a PoR scheme considers the case where an honest client wants to store its file(s) on a potentially malicious server, i.e active adversary. It is a challenge-response interactive protocol where the server proves to the client that its file is intact and retrievable. A PoR scheme comprises five algorithms:

- $\texttt{Setup}(1^\lambda) \rightarrow sk$: a probabilistic algorithm, run by a client, that takes an input a security: 1^λ and outputs a secret key.

- $\texttt{Store}(sk, F) \rightarrow (F^*, \sigma)$: a probabilistic algorithm, run by a client, that takes an input the secret key: sk and a file: F. It encodes F, denoted by F^* as well as generating a set of tags: σ, where F^* and σ is stored on the server.

- $\texttt{GenChal}(|F^*|, 1^\lambda) \rightarrow \vec{c}$: a probabilistic algorithm, run by a client, that takes an input the encoded file size: $|F^*|$ and security: 1^λ. It outputs a set of pairs, $\ddot{c}_j$: (x_j, y_j), where each pair includes a file block index: x_j and coefficient: y_j, both of them are picked uniformly at random.

- $\texttt{Prove}(F^*, \sigma, \vec{c}) \rightarrow \pi$: takes the encoded file: F^*, (a subset of) tags: σ, and a vector of unpredictable random challenges: $\vec{c}$ as inputs and outputs a proof of the file retrievability. It is run by a server.

- $\texttt{Verify}(sk, \vec{c}, \pi) \rightarrow \{0, 1\}$: takes the secret key: sk, vector of random challenges: $\vec{c}$, and the proof π as inputs. It outputs either 0 if it rejects, or 1 if it accepts the proof. It is run by a client.

Informally, a PoR scheme has two main properties: correctness and soundness. The correctness requires that for any key, all files, the verification algorithm accepts a proof generated by an honest verifier. The soundness requires that if a prover convinces the verifier (with a high probability) then the file is stored by the prover; This is formalized

via the notion of an extractor algorithm, that is able to extract the file in interaction with the adversary using a polynomial number of rounds. In contrast to the definition in [46] where GenChal(.) is implicit, in the above we have explicitly defined it, as its modified version plays an integral role in SO-PoR definition (and protocol).

H SO-PoR Model

In this section, we provide a formal definition of SO-PoR. As previously stated, it builds upon the traditional PoR model [46], presented in Appendix G. In SO-PoR, unlike the traditional PoR, a client may not be available every time verification is needed. Therefore, it wants to delegate a set of verifications that it cannot carry out itself. In this setting, it (in addition to file retrievability) must have three guarantees: (a) *verification correctness*: every verification is performed honestly, so the client can rely on the verification result without the need to re-do it, (b) *real-time detection*: the client is notified in almost real-time when server's proof is rejected, and (c) *fair payment*: in every verification, the server is paid only if the server's proof is accepted. In SO-PoR, three parties are involved: an honest client, potentially malicious server and a standard smart contract. SO-PoR also allows a client to perform the verification itself, analogous to the traditional PoR, when it is available.

Definition 10. *A Smart Outsourced PoR (SO-PoR) scheme consists of seven algorithms (*Setup, Store, SolvPuz, GenChall, Prove, Verify, Pay*) defined below:*

- Setup$(1^\lambda, \Delta, z) \rightarrow (\hat{sk}, \hat{pk})$: *a probabilistic algorithm, run by a client. It takes as input a security: 1^λ, time parameter: Δ, and the number of verification delegated: z. It outputs a set of secret and public keys.*

- Store$(\hat{sk}, \hat{pk}, F, z) \rightarrow (\boldsymbol{F}, \sigma, \overrightarrow{\boldsymbol{o}}, aux)$: *a probabilistic algorithm, run only once by a client. It takes as input the secret key: $\hat{sk}$, public key: $\hat{pk}$, a file: F, and the number of verifications: z that the client wants to delegate. It outputs an encoded file: $\boldsymbol{F}$, a set of tags: σ, a set of z puzzles: $\overrightarrow{\boldsymbol{o}}$, and public auxiliary data: aux. First three outputs are stored on the server and last output: aux, is stored on a smart contract.*

- SolvPuz$(\hat{pk}, \overrightarrow{\boldsymbol{o}}) \rightarrow \overrightarrow{\boldsymbol{s}}$: *a deterministic algorithm that takes as input the public key: $\hat{pk}$ and puzzle vector: $\overrightarrow{\boldsymbol{o}}$. It for each j-th verification outputs a pair: $\ddot{s}_j : (v_j, l_j)$ of solutions, where v_j and l_j are outputted at time t_j and t'_j respectively and $t'_j > t_j$. Therefore, the algorithm in total outputs z pairs. Value l_j is sent to the smart contract right after it is discovered. This algorithm is run by the server.*

- GenChall$(j, |\boldsymbol{F}|, 1^\lambda, \ddot{s}_j, aux) \rightarrow \overrightarrow{\boldsymbol{c}}$: *a probabilistic algorithm that takes as input a verification index: j, the encoded file size: $|\boldsymbol{F}|$, security parameter: 1^λ, first component of the related solution pair, $v_j \in \ddot{s}_j$, and public parameters: $pp \in aux$ containing a blockchain and its parameters. It outputs pairs $\ddot{c}_j : (x_j, y_j)$, where each pair includes a pseudorandom block's index: x_j and random coefficient: y_j.*

Also, values x_j are derived from v_j while y_j are derived from pp. This algorithm is run by the server for each verification.

- $\texttt{Prove}(j, \boldsymbol{F}, \sigma, \vec{\boldsymbol{c}}) \to \pi$: *a probabilistic algorithm that takes the verification index j, encoded file: $\boldsymbol{F}$, (a subset of) tags: σ, and a vector of unpredictable challenges: $\vec{\boldsymbol{c}}$, as inputs and outputs a proof of file retrievability. It is run by the server for each verification.*

- $\texttt{Verify}(j, \pi, \ddot{s}_j, aux) \to d : \{0, 1\}$: *a deterministic algorithm that takes the verification index j, proof: π, second component of the related solution pair: $l_j \in \ddot{s}_j$, and public auxiliary data: aux. If the proof is accepted, it outputs $d = 1$; otherwise, outputs $d = 0$. The default value of d is 0. This algorithm is run by the smart contract for each verification and invoked only once for each verification by only the server.*

- $\texttt{Pay}(j, d) \to d' = \{0, 1\}$: *a deterministic algorithm that takes the verification index j, the verification output: d. If $d = 1$, it transfers e amounts to the server and outputs 1. Otherwise, it does not transfer anything, and outputs 0. The default value of d' is 0. The algorithm is run by the contract, and invoked only by $\texttt{Verify}(.)$.*

A SO-PoR scheme must satisfy two main properties: *correctness* and *soundness*. The correctness requires, for any: file, public-private key pairs, and puzzle solutions, both the verification and pay algorithms, i.e. $\texttt{Verify}(.)$ and $\texttt{Pay}(.)$, output 1 when interacting with the prover, verifier, and client all of which are honest. The soundness however is split into four properties: extractability, verification correctness, real-time detection, and fair payment, formally defined below. Before we define the first property, extractability, we provide the following experiment between an environment: $\mathcal{E}$ and adversary: $\mathcal{A}$ who corrupts $C \subsetneq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_\beta\}$, where β is the maximum number of miners which can be corrupted in a secure blockchain. In this game, $\mathcal{A}$ plays the role of corrupt parties and $\mathcal{E}$ simulating an honest party's role.

1. $\mathcal{E}$ executes $\texttt{Setup}(.)$ algorithm and provides public key: $\hat{pk}$, to $\mathcal{A}$.
2. $\mathcal{A}$ can pick arbitrary file F', and uses it to make queries to $\mathcal{E}$ to run: $\texttt{Store}(\hat{sk}, \hat{pk}, F', z) \to (F'^*, \sigma, \vec{\boldsymbol{o}}, aux)$ and return the output to $\mathcal{A}$. Also, upon receiving the output of $\texttt{Store}()$, $\mathcal{A}$ can locally run algorithms: $\texttt{SolvPuz}(\hat{pk}, \vec{\boldsymbol{o}})$ and $\texttt{GenChall}(j, |F'^*|, 1^\lambda, \ddot{s}_j, aux) \to \vec{\boldsymbol{c}}$ as well as $\texttt{Prove}(j, F^*, \sigma, \vec{\boldsymbol{c}}) \to \pi$, to get their outputs as well.
3. $\mathcal{A}$ can request $\mathcal{E}$ the execution of $\texttt{Verify}(j, \pi, \ddot{s}_j, aux)$ for any F' used to query $\texttt{Store}()$. Accordingly, $\mathcal{E}$ informs $\mathcal{A}$ about the verification output. The adversary can send a polynomial number of queries to $\mathcal{E}$. Finally, $\mathcal{A}$ outputs the description of a prover: $\mathcal{A}'$ for any file it has already chosen above.

It is said a cheating prover: $\mathcal{A}'$ is ϵ-admissible if it convincingly answers ϵ fraction of verification challenges [46]. Informally, a SO-PoR scheme supports extractability, if there is an extractor algorithm: $\texttt{Ext}(\hat{sk}, \hat{pk}, P')$, that takes the secret-public keys and the description of the machine implementing the prover's role: $\mathcal{A}'$ and outputs the file: F'.

The extractor can reset the adversary to the beginning of the challenge phase and repeat this step polynomially many times for of extraction, i.e. the extractor can rewind it.

Definition 11 (ϵ-extractable). *A SO-PoR scheme is ϵ-extractable if for every adversary: $\mathcal{A}$ who corrupts $C \subsetneq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_\beta\}$, plays the experiment above, and outputs an ϵ-admissible cheating prover: $\mathcal{A}'$ for a file F', there exists an extraction algorithm that recovers F' from $\mathcal{A}'$, given honest parties public-private keys and $\mathcal{A}'$, i.e. $\text{Ext}(\hat{sk}, \hat{pk}, \mathcal{A}') \to F'$, except with a negligible probability.*

In the above game, the environment, acting on honest parties' behalf, performs the verification correctly; which is not always the case in SO-PoR. As the verification can be run by miners a subset of which are potentially corrupted. Even in this case, the verification correctness must hold, e.g. if a corrupt server sends an invalid proof then even if $\beta - 1$ miners are corrupt (and colluding with it) the verification function will not output 1 and if the server is honest and submits a valid proof then the verification function does not output 0 even if β miners are corrupt, except with a negligible probability. This is formalised below.

Definition 12 (Verification Correctness). *Let β be the maximum number of miners that can be corrupted in a secure blockchain network and λ' be the blockchain security parameter. Also, let $\mathcal{A}$ be the adversary who (plays the above game and) corrupts parties in either $C \subseteq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_{\beta-1}\}$ or $C' \subseteq \{\mathcal{M}_1, ..., \mathcal{M}_\beta\}$. In SO-PoR, we say the correctness of j-th verification is guaranteed if:*

$$\text{in the former case}: Pr[\text{Verify}_C(j, \pi, \ddot{s}_j, aux) = 1] \leq \mu(\lambda')$$
$$\text{in the latter case}: Pr[\text{Verify}_{C'}(j, \pi, \ddot{s}_j, aux) = 0] \leq \mu(\lambda')$$

where $\mu(.)$ is a negligible function.

Also, a client needs to have a guarantee that for each verification it can get a correct result within a (fixed) time period.

Definition 13 (Υ-real-time Detection). *Let $\mathcal{A}$, as defined above, be the adversary who corrupts either C or C'. A client, for each j-th delegated verification, will get a correct output of $\text{Verify}(.)$, by means of reading a blockchain, within time window Υ, after the time when the server is supposed to send its proof to the blockchain network. Formally,*

$$\text{Read}(\Upsilon, \text{Verify}_D(j, \pi, \ddot{s}_j, aux)) \to \{0, 1\}$$

where $D \subsetneq \{C, C'\}$, except with a negligible probability.

Definition 14 (Fair Payment). *SO-PoR supports a fair payment if the client and server fairness are satisfied:*

- *__Client Fairness__: An honest client is guaranteed that it only pays (e coins) if the server provides an accepting proof, except with a negligible probability.*
- *__Server Fairness__: An honest server is guaranteed that the client gets a correct proof if the client pays (e coins), except with a negligible probability.*

Formally, let $\mathcal{A}$ be the adversary who corrupts either C or C', as defined above. To satisfy a fair payment:

$$Pr[\texttt{Pay}_D(.) = b \cap \texttt{Verify}_D(.) = b] \geq 1 - \mu(\lambda'), \qquad (2)$$

the following inequality must hold:

$$Pr[\texttt{Pay}_D(.) = b' \cap \texttt{Verify}_D(.) = b] \leq \mu(\lambda'), \qquad (3)$$

where $D \subsetneq \{C, C'\}, b \neq b'$, and $b, b' \subsetneq \{0, 1\}$

The above definition also takes into account the fact that the client at the time of delegated verification is not necessarily available to make the payment itself, so the payment is delegated to a third party, e.g. a smart contract. In this case, the definition ensures that even if the client or/and server are honest, the third party cannot affect the fairness (except with a negligible probability).

Definition 15 (SO-PoR Security). *A SO-PoR scheme is secure if it is ϵ-extractable, and satisfies verification correctness, Υ-real-time detection, and fair payment properties.*

Remark 9. The folklore assumption is that (in a secure blockchain) a smart contract function *always outputs a correct result*. However, this is not the case and it may fail under certain circumstances. For instance, as shown in [36] all rational miners may not verify a certain transaction. As another example, an adversary (although with a negligibly small probability) discards a certain honestly generated blocks, reverses the state of blockchain and contract, or breaks a client's signature scheme. Accordingly, in our definitions above, we take such cases into consideration and allow the possibility that a function outputs an incorrect result even though with a negligibly small probability.

Remark 10. SO-PoR model differs from traditional (e.g. [26, 46]) and outsourced PoR (e.g. [3, 53]) models in several aspects. Only the SO-PoR model offers all the properties. In particular, traditional PoRs only offer extractability while outsourced ones offer liability as well, that allows a client (by re-running all verifications function) to detect a verifier if it provides an incorrect verification output, so the client cannot rely on the verification result provided. As another difference, the SO-PoR model takes into account the case where an adversary can corrupt both the server and some miners at the same time.

Remark 11. SO-PoR should also support the traditional PoR where only client and server interact with each other (e.g. client generates challenges, and verifies proof) when the client is available. To let SO-PoR definition support that too, we can simply define a flag: ξ, in each function, such that when $\xi = 1$, it acts as the traditional PoR; otherwise (when $\xi = 0$), it performs as a delegated one. For the sake of simplicity, we let the flag be implicit in the definitions above, where the default is $\xi = 1$

I   SO-PoR Security Proof

In this section, we first restate the main security theorem of SO-PoR protocol (presented in Section 5.3) and then prove it.

Theorem 2. *SO-PoR protocol is secure (according to Definition 15) if the tags/MAC's are unforgeable, $\mathrm{PRF}(.)$ is a secure pseudorandom function, the blockchain is secure, C-TLP protocol is secure, and $\mathrm{H}(\mathcal{B}_\gamma||...||\mathcal{B}_\zeta)$ outputs an unpredictable random value (where $\zeta - \gamma$ is a security parameter).*

Proof (sketch). In the following, we prove that SO-PoR protocol satisfies every property that were defined in Appendix H.

Verification correctness (according to Definition 12). We first argue that the adversary who corrupts either $C \subseteq \{\mathcal{M}_1, ..., \mathcal{M}_\beta\}$ or $C' \subseteq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_{\beta-1}\}$ with a high probability, cannot influence the output of $\mathrm{Verify}(.)$ performed by a smart contract in a blockchain; in other words, the verification correctness holds. In short, the verification correctness boils down to the security of the underlying blockchain. In the case where the adversary corrupts C (when the server provides an accepting proof), for the adversary to make the verification function output 0, it has to: (a) either forge the server's signature, or (b) fork the blockchain so the chain comprising the accepting proof is discarded. In case (a), if it manages to forge the signature, it can generate a transaction that includes a rejecting proof where the transaction is signed on the server's behalf. In this case, it can broadcast the transaction as soon as the transaction containing an accepting proof is broadcast, to make the latter transaction stale. Nevertheless, the probability of such an event is negligible, $\epsilon(\lambda')$, as long as the signature is secure. Moreover, due to the liveness property of blockchain, an honestly generated transaction will eventually appear on an honest miner's chain [22]. In case (b), the adversary has to generate long enough (valid) chain that excludes the accepting proof, but this also has a negligible success probability, $\epsilon(\lambda')$, under the assumption that the hash power of the adversary is lower than those of honest miners (i.e. under the honest majority assumption) and due to the liveness property. Now we turn our attention to the case where the adversary corrupts C' (when the server provides a rejecting proof). In this case, for the adversary to make the verification function to output 1, the honest miners must not validate the transaction that contains the proof. Nevertheless, as long as the blockchain is secure and the computational advantage of skipping transaction validation is low, i.e. the validation imposes a low computation cost, the miners check the transaction's validation [36]. Also, as shown in [36] when a transaction validation imposes a high computation cost, two generic techniques can be used to support exact or probabilistic correctness (of a smart contract function output). We conclude the correctness of $\mathrm{Verify}(.)$ output is guaranteed with a high probability.

ϵ-extractable (according to Definition 11). In the following, we show that if a proof produced by an adversary: $\mathcal{A}$ who corrupts $C' \subseteq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_{\beta-1}\}$ is accepted by $\mathrm{Verify}(.)$ with probability at least ϵ, then the file can be extracted by a means of an extraction algorithm. As mentioned before, $\mathrm{Verify}(.)$ can use both disposable and permanent tags, in the latter case $\mathrm{Verify}(.)$ is run by a smart contract, while in the former one the client runs it. For the sake of simplicity, we first consider the case where $C' = \mathcal{S}$. In this case, the extractability proof is similar to the one in [46], with a few differences, in SO-PoR: (a) the extractor can use both disposable tags and permanent tags (when the former run out), (b) assumes C-TLP protocol is secure, (c) assumes $\mathrm{H}(\mathcal{B}_\gamma||...||\mathcal{B}_\zeta)$ out-

puts a random value even if β miners try to influence its output. Note that in [46] only the permanent tags are used, and since the client generates the challenges and performs the verification, it does not use other primitives; hence, it does not require other security assumptions. As proven in Theorems 4 and 5, C-TLP protocol is secure. Moreover, as analysed and proven in [1, 3], an output of $\mathtt{H}(\mathcal{B}_\gamma||...||\mathcal{B}_\varsigma)$ is random value even if some blocks are generated or selectively disseminated by malicious miners. We conclude that the extractor can extract the file when $C' = \mathcal{S}$, and the extractor is interacting a ϵ-admissible prover. Now we move on to the case where $C' \subseteq \{\mathcal{S}, \mathcal{M}_1, ..., \mathcal{M}_{\beta-1}\}$. In this case, the proof (provided by $\mathcal{S}$) is rejecting but the corrupt miners may try to make $\mathtt{Verify}(.)$ output 1. Note, if they succeed to do so, then the file can be extracted only by using the permanent tags but not the disposable ones. However, as shown above (i.e. due to the verification correctness), they have a negligible probability of success, when the blockchain is secure.

Υ-*real-time Detection* (according to Definition 13). In the following, we argue that after the server broadcasts a proof at a certain time, say t, to the network, the client can get a correct output of $\mathtt{Verify}(.)$ at most after time period Υ, by a means of reading the blockchain. The proof is split into two parts: (a) correctness of $\mathtt{Verify}(.)$ output, and (b) the maximum delay on the client's view of the output. Since we have already shown above that (when C or C' is corrupt) the correctness of $\mathtt{Verify}(.)$ output is guaranteed, we focus on the latter property, i.e. the delay. To describe the delay, we need to recall two blockchain notions: *liveness* and *slackness* [22, 7]. Informally, liveness states that an honestly generated transaction will eventually be included more than k blocks deep in an honest party's blockchain [22]. It is parameterised by wait time: u and depth: k. We can fix the parameters as follows. We set k as the minimum depth of a block considered as the blockchain's *state* (i.e. a part of the blockchain that remains unchanged with a high probability, e.g. $\mathtt{k} \geq 6$) and u the waiting time that the transaction gets k blocks deep. As shown in [7], there is a slackness on honest parties' view of the blockchain. In particular, there is no guarantee that at any given time, all honest miners have the same view of the blockchain, or even the state. But, there is an upper-bound on the slackness, denoted by $\mathtt{WindowSize}$, after which all honest parties would have the same view on a certain part of the blockchain state. This means when an honest party (e.g. the server) propagates its transaction (containing the proof) all honest parties will see it on their chain after at most: $\Upsilon = \mathtt{WindowSize} + \mathtt{u}$ time period. So when the adversary corrupts C or C', but in the latter case the server constructs a valid transaction (regardless of the proof status) the client by reading the blockchain (i.e. probing the miners) can get a correct result after at most time period Υ when the server sends the proof. Also, when parties in C' are corrupt and the transaction (containing the proof) is not valid, as discussed above (for verification correctness) the honest miners would detect the invalid transaction and do not include in the chain, therefore the output of $\mathtt{Verify}(.)$ would be the same as its default value: 0; the same holds when the server sends nothing to the network. This concludes the proof related to Υ-real-time detection in SO-PoR protocol.

Fair Payment (according to Definition 14). The proof takes into consideration that the correctness of an output of $\mathtt{Verify}(.)$ is guaranteed (as shown above). It boils down to

the correctness of `Pay(.)` as it is interrelated to the output of `Verify(.)`. In particular, for the adversary to make inequality 3 not hold, it has to break `Pay(.)` correctness, e.g. pays the server despite the verification function outputs 0. Therefore, it would suffice to show the adversary who corrupts C or C' cannot affect `Pay(.)` output's correctness. To do so, we can apply the same argument used to prove the correctness of `Verify(.)` above, as the correctness of `Pay(.)` relies on the security of the blockchain as well. □

J Further Remarks on SO-PoR Protocol

Remark 12. In each verification, e.g. j-th one, it is required that the server can: (a) learn the random challenges, (b) compute a proof, and (c) record it in the smart contract, before it is able to learn key l_j; otherwise, (i.e. if it learns l_j before sending and registering the proof), it can tamper with the data and pass the verification. Because by knowing l_j it can construct valid tags for the data that has been tampered with. That is why, in the protocol, it is required: $t'_j \geq t_j + \Delta_1 + \Delta_2$

Remark 13. The way disposable tags are generated in SO-PoR differs from those computed in traditional/outsourced PoR schemes, in spite of having similarities structure-wise. Specifically, (unlike existing protocols, e.g. [46, 3]) in SO-PoR, each random value, $r_{b,j}$, utilised to generate a disposable tag of a block (for j-th verification), is not derived from the block index. Instead, it depends on (a fresh secret key for j-th verification and) the total number of blocks challenged in each verification that is a public value. This means, the verifier does not need to know and verify each challenged block's index in the verification phase, which leads to a lower cost.

Remark 14. With minor adjustments, we can reduce the smart contract storage cost from $O(z)$ to constant, $O(1)$ and offload the cost to the server. The idea is that the client after computing the commitmnet vector: $\vec{h} = [h_1, ...h_z]$, in step 2d, it preserves the ordering of the elements (i.e. h_j is associated with j-th verification) and constructs a Merkle tree on top of them. It stores the tree and the vector on the server, and stores only the tree's root: R, on the contract. In this case, the server in step 3e after recovering $\ddot{p}_j = (l_j, d_j)$, computes: $h_j = \mathtt{H}(l_j||d_j)$, and sends a Merkle tree proof (that h_j corresponds to R) along with $\ddot{p}_j$ to the contract. In step 4b, the contract: (a) checks if $h_j = \mathtt{H}(l_j||d_j)$, and (b) verifies the Merkle tree proof. The rest remains unchanged. As a result, the number of values stored in the contract is now $O(1)$. This adjustment comes with an added communication cost: $O(|h_j| \log z)$ for each verification. Nevertheless, the added cost is small and independent of the file size. For instance, when $z = 10^6$ and $|h_j| = 256$, the added communication cost is only about 5.1 kilobit.

Remark 15. One might be willing to use a combination of existing publicly verifiable PoR and smart contract, such that the contract performs the verification on the client's behalf. However, this approach would have a higher computation or communication cost (especially in the verification phase) than our protocol. Specifically, there exist two publicly verifiable PoR schemes, based on either (a) BLS signatures, e.g. [46], or (b) Merkle tree, e.g. [40]. The former approach, in total, requires zc exponentiations in the verification phase, whereas SO-PoR requires no exponentiations in this phase. Also,

the BLS signature-based scheme takes a very long time to encode a file, as the required number of exponentiations is $O(|\boldsymbol{F}|)$. For instance, as measured in [3], it takes about 55 minutes to encode a 64-MB file, meaning it would take about 14 hours to encode a 1-GB file. But, in SO-PoR the number of exponentiations, in the store phase, is independent of and much fewer than the file size. On the other hand, the proof size in the Merkle tree-based approach is logarithmic with the file size, i.e. $256zc\log|\boldsymbol{F}|$ bits, that leads to a high server-side's communication cost. By contrast, the proof size in SO-PoR is independent of the file size and is much shorter, i.e. $884z$ bits.

K   SO-PoR's Full Evaluation

In the following, we provide a full analysis of SO-PoR and compare its properties and detailed costs to those protocols that support outsourced PoR (O-PoR), i.e. [3, 53, 5]. Note, there are two protocols proposed in [5]; in the following, we only consider the one that supports public verifiability, i.e. basic PoSt. Recall, we consider a generic case where a client outsources z verifications and in our cost analysis, we also compare SO-PoR cost with the cost of the most efficient privately verifiable PoR [46], too.
Properties. We start with a crucial feature that any O-PoR must have: real-time detection. Recall, real-time detection requires a client to receive a correct verification result in (almost) real-time without the need for it to re-execute the verification itself. This is offered only by SO-PoR. By contrast, in [3] the auditor may never notify the client, even if it does, its notification would not be reliable, and the client has to redo the verification to verify the auditor's claim. Similarly, in [53] the client has to fully trust the auditor to get notified on-time. The basic PoSt in [5] requires the server to collect all PoR's and send them to a validator in one go. This means, the client and validator cannot detect data tampering in real-time if a subset of the PoR's are invalid; instead they need to wait until z PoR's are collected by the server. So, [3, 53, 5] are not suitable for the cases where a client must be notified by a potentially malicious auditor as soon as an unauthorised modification on the sensitive data is detected. The fair payment is another vital property in O-PoR, as the cloud server and auditor must be paid fairly, in the *real world* when they serve a client. This feature is explicitly captured by only SO-PoR. The protocols in [3, 53] do not have any mechanism in place. In [3], one may allow the auditor to pay the server on the client's behalf. But, this is problematic. The server and auditor can collude to save costs, in a way that the server generates accepting proofs for the client but generates no proof for the auditor, and still the auditor pays it. This violates the fair payment and cannot be detected by the client unless it performs all the verification itself. On the other hand, a client in [53] has to fully trust the auditor (with the payment too), otherwise the auditor can collude with the server to violate the fair payment. The authors of [5] briefly state that the basic PoSt's verification can be performed by a smart contract who, after ensuring the proofs are valid, pays the server. Nevertheless, as stated previously, the verification cost of this protocol is too high for a smart contract, i.e. imposes at least z modular exponentiations over RSA modulus and requires a high number of messages logarithmic with the file size to be sent to the contract. Thus, even though it can support fair payment *in theory*, it is very costly in practice. Another property is the cost of onboarding a new verifier, as it determines how

flexible the client can be, to pick a new auditor when its current one is misbehaving. This cost in [3] is significantly high, as it requires the verifier to download the entire file, generate metadata, and prove in zero-knowledge the correctness of metadata to the client. But, that cost in SO-PoR and [53, 5] is very low as the client only sends them a small set of parameters without the need to access the outsourced data. Furthermore, as stated above, a client in [53] has to fully trust the auditor with the correctness of verification but this is not the case in SO-PoR and [3, 5], as they consider a potentially malicious auditor (under different assumptions). Table 5 summaries the result of the comparison between the four protocols' main properties.

Table 5: O-PoR Property Comparison. $\checkmark^*$ indicates the property is met only in theory.

Protocols	Properties		
	Real-time Detection	Fair Payment	Untrusted Auditor
SO-PoR	$\checkmark$	$\checkmark$	$\checkmark$
[3]	$\times$	$\times$	$\checkmark$
[53]	$\times$	$\times$	$\times$
[5]	$\times$	$\checkmark^*$	$\checkmark^*$

Computation Complexity. In our analysis, we do not take into account the cost of erasure-coding a file, as it is identical in all schemes. We first analyse the computation cost of SO-PoR. A client in step 2b performs n multiplications and n additions to generate permanent tags. In step 2c, it performs cz multiplications and cz additions to generate disposable tags for z verifications. The client in step 2d invokes GenPuz(.) function, in C-TLP, that costs $O(z)$. So, the client's total computation cost of preparing and storing a file is $O(n + cz)$. Now we consider the cloud server's cost that can be categorised into two classes: (a) solving a puzzle: SolvPuz(), run only once, and (b) generating PoR run for each verification. In particular, the cloud in step 3a, invokes SolvPuz(), in C-TLP, that costs $O(Tz)$ this includes the cost in step 3e as well. To compute proofs, in step 3c, it performs $2cz$ multiplications and $2cz$ additions. So, the server to generate proof is $O(cz)$. Next, we analyse the cost of the smart contract. In step 4b, it invokes Verify(.), in C-TLP, that in total costs $O(z)$, this involves invoking z hash function's instances. Also, the contract in step 4c performs $z(1+c)$ and $z(1+c)$ modular multiplications and additions respectively. Thus, the total cost is $O(cz)$ involving mainly modular additions and multiplications.

Now we analyse the computation cost of [3]. To prepare file tags, a client performs: n multiplications and n additions. Also, to verify tags generated by the auditor, the client has to engage in a zero-knowledge protocol that requires it to carry out $6n$ exponentiations and $2n$ multiplications. Therefore, the client's computation complexity is $O(n)$. Also, the auditor in total performs $3n$ multiplications, $3n$ additions and $3n$ exponentiations to prepare file's metadata, so its complexity at this phase is $O(n)$. For the cloud to generate z proofs, in total it performs $2z(c + c')$ multiplications and the same number of additions, where c' is the number of challenges sent by the auditor to the cloud (on client's behalf) and $c > c'$, e.g. $c' = (0.1)c$. So, the total complexity of

the cloud is $O(z(c + c'))$. Next, we consider the verification cost. The auditor performs $z(1 + c)$ multiplications and $z(1 + c)$ additions to verify PoR. It also performs $2cz$ additions in *CheckLog* algorithm that requires the client to perform $z(2c + 1)$ additions and cz multiplications. Nevertheless, as discussed in Section 2, running only *CheckLog* does not allow the client to detect a misbehaving auditor. Thus, it has to run *ProveLog* too, that requires the client to perform cz multiplications and cz additions and requires the auditor to reveal all its secrets to the client. So, the total verification complexity is $O(cz)$. Now we turn our attention to the computation complexity of [53]. To prepare metadata, the client needs to perform $2n$ multiplication and $2n$ additions, so the client's complexity is $O(n)$. For the cloud to generate a proof it needs to perform $3z$ exponentiations, $z(3c + 6)$ multiplications and $z(3c + 2)$ additions. So, its complexity is $O(cz)$. On the other hand, the verifier performs cz exponentiations to compute challenges. To verify the proof, it carries out $6z$ exponentiations, cz multiplications, cz additions, and $7z$ pairings. Therefore, the verifier complexity is $O(cz)$ dominated by expensive exponentiations and pairing operations. Also, we analyse the computation cost of efficient privately verifiable PoR in [46]. A client in the store phase performs n multiplications and n additions to construct the tags. So, its complexity is $O(n)$. A server performs $2cz$ multiplications and $2cz$ additions to generate proofs, so in total $4cz$ or $O(cz)$ modular operations for z verifications it carries out. The client, as verifier this time, performs in total $2z(1 + c)$ or $O(z(1 + c))$ modular operations. Next, we turn our attention to [5], and evaluate the cost of the protocol that supports public verifiability, basic PoSt. As stated by the authors, they add a VDF to the PoR construction in [35] which uses a Merkle tree-based PoR. Since, this PoR scheme only involves invocations of hash function, here we only focus on VDF cost as it dominates other computation costs. We assume that the most efficient publicly verifiable delay function VDF [51] is used. In the setup, the client constructs a Merkle tree on the entire data by invoking a hash function many times and also generates a random challenge. We ignore these costs as they are dominated by VDF's costs. To generate z PoR's, the server invokes VDF z times that in total runs in time period T. This involves $3Tz$ modular exponentiations over $\mathbb{Z}_N$ (where N is a RSA modulus), and Tz modular multiplications. Therefore, its complexity is $O(Tz)$. For a validator to ckeck the proofs output by VDF, it performs $3z$ modular exponentiations over $\mathbb{Z}_N^*$ (or $\mod \phi(N)$). So the verifier's complexity is $O(z)$.

Now we compare the protocols above. The verification in SO-PoR is much faster than the other three protocols; firstly, it requires no exponentiations in this phase, whereas [53, 5] do, and secondly, it requires $\frac{9c+3}{2(1+c)}$ times fewer computation than [3]; Specifically, when $c = 460$, SO-PoR verification requires about 4.5 times fewer computation than the verification in [3] needs. In SO-PoR, the cloud server needs to perform Tz exponentiations to solve puzzles, however this is independent of the file size. The sever in [5] also performs $3Tz$ modular exponentiations, which is 3 times higher than the number of exponentiations done by the server in SO-PoR. The protocols in [3, 46] do not include the puzzle-solving procedure (and they do not offer all features that SO-PoR does). Furthermore, the proving cost in SO-PoR is similar to that of in [3], and is much better than [53], as the latter one requires both exponentiations and pairing operations while the prove algorithm in SO-PoR does not involve any exponentiations. The proving cost in [5] is the lowest, as it requires only invocations of a hash function.

Also, the store phase in SO-PoR has a much lower computation cost than the one in [3]. The reason is that the number of exponentiations required (in this phase) in SO-PoR is independent of file size and is only linear with the number of delegated verifications; however, the number of exponentiations in [3] is linear with the file size. For instance, when $||\boldsymbol{F}|| = 1$-GB, the total number of blocks is: $n = \frac{1-\text{GB}}{128-\text{bit}} = 625 \times 10^5$. Since the number of exponentiations in [3] is linear with the number of blocks, i.e. $9n$, the total number of exponentiations imposed by store algorithm is: 5625×10^5 which is very high. This is the reason why in the experiment in [3] only a small file size: 64-MB, is used, that can be stored locally without the need to use cloud storage, in the first place. Now, we turn our attention to SO-PoR. Let the verification be done every month for a 10-year period, in this case, $z = 120$. So, the total number of exponentiations required by the store in SO-PoR is 121. This means the store phase in SO-PoR requires over 46×10^5 times fewer exponentiations than the one in [3] needs. On the other hand, the store algorithm in [53] does not involve any exponentiations; however, its number of modular additions and multiplication is higher than the ones imposed by the store in SO-PoR. The store cost in [5] is the lowest, as it requires only invocations of a hash function. Furthermore, the verification and prove cost of SO-PoR and privately verifiable PoR [46] are identical.

Communication Complexity. In our analysis, we do not take into account the communication cost of uploading an encoded file, i.e. $||\boldsymbol{F}||$, when the client for the first time sends it to the cloud, as it is identical in all schemes. The communication cost of SO-PoR is as follows. The client, in step 2e, sends n permanent tags, zc disposable tags, and the output of $\texttt{GenPuz}(.)$ to the server and contract, where each (permanent/disposable) tag: $\sigma_j \in \mathbb{F}_p$ and $|\sigma_j| = 128$-bit. Note the client also sends a few public parameters: $\hat{pk}$, whose size is short. Therefore, the client's bandwidth is: $128(n + cz + 19z)$ bits, while its communication complexity is $O(n + cz)$. The cloud in step 3d, sends z pairs (μ_j, ξ_j), where $\mu_j, \xi_j \in \mathbb{F}_p$ and $|\mu_j| = |\xi_j| = 128$-bit. Also, in step 3e, it sends to the contract the output of $\texttt{Prove}(.)$, in C-TLP, whose total size is $628z$ bits. So, the clouds total bandwidth is about $884z$ bits and its complexity is $O(z)$ which is independent of and constant in the file size.

The communication cost of [3] is as follows. The client sends n tags to the server, where the size of each tag is about 128 bits. So its bandwidth is $128n$, and its complexity is $O(n)$. The auditor also sends n tags to the server, where each tag size is also 128 bits. It also sends the tags to the client along with zk proofs that contain $4n$ elements in total, where $2n$ of them are elements of $\mathbb{Z}_N$ and each element size is 2048 bits, and each of the other $2n$ elements is 160 bits long. Also, the auditor in *ProveLog* sends z pairs to the client with the bandwidth of $256z$. So, the auditor's total bandwidth and complexity is $4672n + 256z$ and $O(n + z)$ respectively. Moreover, the cloud sends the entire file, F, to the auditor in the store phase and also sends $2z$ pairs of PoR to the auditor, where each element of the pair is of size 128 bits. Therefore, the cloud's total bandwidth is $||\boldsymbol{F}|| + 256z$, while its complexity is $O(||\boldsymbol{F}|| + z)$. Note that the cloud's proof size complexity is constant and independent of the file size, i.e. $O(1)$. Now, we analyse the communication cost of [53]. The client bandwidth and complexity are $2048n$ and $O(n)$ respectively, as it sends to the cloud $2n$ tags, where each tag size is

1024 bits. Also, the cloud bandwidth and complexity are $6144z$ and $O(z)$ respectively, as for each verification the cloud sends to the verifier 6 elements each of them is 1024-bit long. Furthermore, in [46] the client bandwidth in the store phase is $128n$, while the server bandwidth is $256z$. In this scheme, the complexity of a proof size is $O(1)$. In [5], the client can send only the file and random challenge of size 128-bit to the server, who creates a Merkle tree on top of the file's blocks. Therefore, the client's bandwidth is 128-bit. The server sends to the verifier cz PoR proofs that cost it in total at least $128cz\log(n)$ bits. Also, the server sends VDF's proofs for z outputs, that cost in total $4096z$. Therefore, the server's total bandwidth is $(128cz\log n) + 4096z$. The reason the cost involves c (the number of challenges) is that unlike the other three schemes, this scheme does not support a linear combination of tags/proofs (e.g. homomorphic tags), and in each proving phase c proofs are generated. Furthermore, in this scheme, the complexity of a proof size is logarithmic with the number of file blocks, $O(\log(n))$.

To conclude, the verifier-side bandwidth of SO-PoR (and [53, 5]) is much lower than [3]. For instance, when $||\boldsymbol{F}|| = 1$-GB and $z = 100$, a verifier in SO-PoR requires 62×10^6 fewer bits than the one in [3] does. A client in SO-PoR has a higher bandwidth than it would have in the rest of the protocols. But, this cost is one-off, at the setup phase. The server-side bandwidth of SO-PoR is the lowest; for instance (for the same parameters above) a server in SO-PoR requires 9×10^4, 7, and 1729 times fewer bits than those required in [3], [53] and [5] respectively. Moreover, [5] has the worst proof size complexity, which is logarithmic to the file size; while the proof size complexity of the rest of the schemes is constant. Thus, SO-PoR's server-side bandwidth is significantly lower than the rest while having constant proof size.

Remark 16. In [3], the additional costs to secure parties against a malicious client stem from only the store phase, where an auditor downloads the entire file, generates zero-knowledge proofs, and has the client sign them after verifying the proofs. Therefore, the overheads of proving and verifying phases, in this protocol, would remain unchanged if the protocol considers an honest client.

L More Efficient Proof of Storage-Time

In the following, we show how C-TLP can be used in the PoSt protocols that were proposed in [5] to improve their costs. As stated previously, two protocols: basic PoSt and compact PoSt, supporting proof of storage-time (in the random oracle) are proposed in [5], where basic PoSt uses VDF and is publicly verifiable while compact PoSt utilises a trapdoor delay function (TDF) and is privately verifiable. Also, recall that VDF/TDF is used to allow the server to derive multiple challenges at different points over a certain time period T. Note, both types of delay function (VDF and TDF) impose the same computation cost to the server, i.e. $3Tz$ modular exponentiations and Tz modular multiplication if the fastest delay function is used [51].

Now we show how to replace the delay function in these schemes with C-TLP, in the random oracle. As in the PoSt protocols, the client at the setup precomputes random challenges (and their PoR tags). But, it encodes all challenges for $z-1$ PoR proofs (excluding first one) into puzzles using C-TLP. It sends to the server all puzzles, encoded

file and plaintext challenges for the first PoR proof. As before, the server generates the first PoR proof using the challenges sent to it in the plaintext. Nevertheless, to find j-th challenge to generate j-th PoR proof, the server solves the related puzzle, where $j > 1$. It sends to the client all PoR proofs (or a combination of them) after period T. Also, if the basic PoSt is used, the server sends C-TLP proofs that can be efficiently verified by anyone. On the other hand, if the compact PoSt is used, then the server does not need to send the C-TLP's proofs, as the client already knows the random challenges. The adjustments considerably improve the PoSt protocols' costs. In particular, the server's computation cost would be $\frac{1}{3}$ of the costs imposed by either of the original PoSt protocols. Also, there will be $3z$ further reduction in the number of exponentiations: (a) at the verifier side, in the basic PoSt, as it does not need to perform any exponentiation to check the correctness of C-TLP's output, (b) at the client-side, in the compact PoSt, as the client does not need to evaluate TDF at the setup to precompute the challenges which in total involves $3z$ modular exponentiations over $\phi(N)$. Moreover, the proof size would be reduced by a factor of 6.5.

Remark 17. Although the use of C-TLP in the PoSt protocols can reduce the computation and communication costs, (a) the server-side I/O cost in these schemes (i.e. $O(\log n)$ in the basic PoSt and $O(n)$ in the compact one), and (b) PoR's proof size complexity in the basic PoSt, will remain the same. Because these costs stem from the underlying PoR schemes, used as a black box, by the two PoSt protocols.